

Burgers, Lipstick & Underwear

What Strange Indicators Really Tell Us About
the Economy

Anton Biletskyi-Volokh

Copyright © Anton Biletskyi-Volokh. All rights reserved.

ISBN, publisher imprint, edition details and pricing to be assigned before final submission.

Contents

- Introduction — The World, Translated Badly but Usefully
- Chapter 1 — The Measure You Can Eat
- Chapter 2 — Prices Become Human
- Chapter 3 — The Folklore of Recession
- Chapter 4 — The Economy Under the Economy
- Chapter 5 — Air You Can Count
- Chapter 6 — When Measures Fight Back
- Chapter 7 — Benchmarks for Machines
- Chapter 8 — How to Read a Weird Index

Introduction — The World, Translated Badly but Usefully

The economy is too large to see.

That is the first problem. It is not hidden in the cinematic sense. It does not sit behind a locked door, waiting for one brave analyst with a torch. It is everywhere at once: in a shipping rate, a wage slip, a supermarket shelf, a mortgage payment, a lipstick counter, a container port, a polluted street, a benchmark leaderboard. It is too distributed, too abstract, too full of moving parts. People cannot hold it in their heads.

So they translate it.

They turn exchange rates into burgers. They turn purchasing power into hours of work. They turn recession anxiety into lipstick and underwear. They turn invisible particles in the air into cigarettes. They turn machine intelligence into leaderboard scores. Sometimes the translation is brilliant. Sometimes it is nonsense wearing a clever hat. Often it is both, which is why it survives.

An official statistic can be more accurate and less useful at the same time. A consumer price index may be methodologically serious, but nobody feels one in the hand. A Big Mac is not a national accounts system, but it is a thing a person can imagine buying. A ship carrying iron ore is not the whole global economy, but a dry bulk freight rate can say something real about physical demand. A benchmark for software-fixing AI does not tell us whether a machine is intelligent in any satisfying philosophical sense, but it does tell us whether a model can survive a very specific kind of practical test.

This book is about those translations.

It is not a catalogue of quirky indicators. A catalogue would be easy and dreadful: one index per page, a little anecdote, a neat moral, repeat until everyone involved loses the will to live. The useful question is not “How many weird indices exist?” The useful question is: why do some strange measures clarify the world while others become attractive folklore?

The answer is not simply accuracy. Accuracy matters, obviously. But strange indicators succeed because they compress. They take a system that feels remote and give it a handle. A burger becomes a currency argument. A phone becomes a wage argument. Lipstick becomes a recession story. Cigarettes become a pollution analogy. The handle is powerful because it is familiar. It is dangerous for exactly the same reason.

The human mind likes a good proxy. It likes one object standing in for a system, one vivid unit replacing an abstraction, one ranking telling us where we are. This habit is not stupid. It is how we navigate a world too complicated for direct perception. The problem begins when the proxy stops being a translation and starts pretending to be the thing itself.

A good strange indicator says: this is a useful way to look.

A bad one says: this is reality.

That difference is the book.

There is a second difference running underneath it: measurement and meaning are not the same thing.

A measurement can be technically correct and still fail to mean much to a human being. A price index can be serious and leave the reader unmoved. A benchmark can be well designed and still be misread by everyone quoting it. A ranking can compress reality efficiently and then be used as if it were reality. Numbers do not become meaningful merely by being numbers. They become meaningful when they connect to a question people understand.

This is why strange indicators are worth studying rather than merely collecting. They reveal the negotiation between technical measurement and human attention. They show how a unit can carry emotion, memory, bias, status and comedy. They show why a rough proxy may change public understanding faster than a perfect table. They also show why people who communicate data need more than accuracy. They need judgment about what the reader will do with the number once it lands.

The book's title objects are deliberately ordinary. Burgers, lipstick and underwear do not sound like respectable economic instruments. That is part of their value. They strip away the ceremonial seriousness around measurement and expose the basic act: choose a thing, compare it across contexts, infer something larger. Sometimes the inference is disciplined. Sometimes it is a magic trick. The reader's job is to learn which is which.

The argument moves from the most accessible proxy to the most modern one. A burger introduces purchasing power. A phone turns price into labor time. Lipstick and underwear show how consumer folklore seduces. Freight reveals a harder physical economy. Cigarette equivalents show how risk becomes readable. Governance cases show how measures change behavior. AI benchmarks prove the same old index problem has moved into the newest machines. The final chapter turns the journey into a way of reading.

This is not a book against weird measures. It is a book against lazy belief in them.

The distinction matters because suspicion can become as stupid as credulity. There is a type of reader who thinks every popular comparison is fake because it simplifies. That reader is not being rigorous; they are often hiding from communication. All explanation simplifies. The question is whether the simplification is honest about what it loses.

The best strange indicators are honest losses. They say: I am not the whole system, but I can show you one useful angle. The worst ones hide the loss and sell the angle as truth.

That is the standard to use with every strange measure in this book.

Chapter 1 — The Measure You Can Eat

The Big Mac is a ridiculous ambassador for international economics, which is probably why it works.

It is not elegant. It does not have the solemn look of an equation. It arrives in a box, sheds lettuce, smells faintly of industrial certainty and can be eaten by a person who has no desire to discuss purchasing power parity over lunch. Yet for decades it has done something most macroeconomic explanations fail to do: it has made exchange-rate theory digestible in under a minute.

The idea is famous because it is almost offensively simple. If the same product is sold in many countries, compare its price across those countries. Convert the local price into a common currency. Ask whether the result looks cheap or expensive relative to the exchange rate. The Economist has published the Big Mac Index since the 1980s and now maintains public data and calculation code for it, including local prices, exchange rates and GDP-adjusted variants. [note 1]

That does not make the burger a perfect model. It makes it a very good door.

Purchasing power parity is one of those ideas that sounds worse than it is. In plain English, it asks whether money buys the same amount of stuff in different places. If a product costs much more in one country than another after exchange-rate conversion, something interesting may be happening: the currency may be overvalued or undervalued, local costs may differ, taxes may bite, wages may be different, competition may be weak, or the product may simply occupy a different place in local culture.

The Big Mac works because it turns that whole mess into a sandwich.

This is not because McDonald's is economically pure. Quite the opposite. A Big Mac is full of local contamination: rents, wages, advertising costs, supply chains, taxes, market strategy and the strange social role of eating American fast food in a given city. That contamination is the point. The burger is not a laboratory object. It is a small commercial animal living in the real world.

If the index were a perfect measurement instrument, it would be less interesting. Its value lies in being rough, memorable and surprisingly revealing. It is close enough to the economic idea to teach it, and familiar enough for people to care.

But the same qualities invite abuse.

The first abuse is treating the burger as a currency oracle. A Big Mac price can suggest that a currency looks cheap or expensive against another. It cannot tell a trader what to do on Tuesday morning. It cannot replace a central bank model, a current-account analysis, a productivity comparison or a serious inflation measure. A burger is not a portfolio strategy. Anyone who says otherwise should probably be kept away from both economics and lunch.

The second abuse is pretending the product is identical everywhere in every relevant sense. It is standardized enough to be useful, but not metaphysically identical. Ingredients, taxes, labor costs, regulation, market positioning and local demand differ. In some countries, McDonald's is everyday food; in others, it is a more expensive international brand. A product can be the same menu item and still carry different social meaning.

The third abuse is forgetting coverage. The Big Mac Index only exists where the product exists and where data can be collected. A global index built around one company's sandwich has obvious holes. Those holes do not destroy it. They define its scope.

This is where strange indicators become interesting. The Big Mac Index is not important because burgers secretly run the global economy. It is important because it demonstrates the minimum viable magic of a good proxy:

- the object is familiar;
- the method is explainable;
- the comparison is portable;
- the limitations are obvious enough to discuss;
- the result invites argument rather than obedience.




Figure 1. Big Mac dollar prices in selected economies, using The Economist data.

That last point matters. Good indicators do not end thought. They start better arguments.

A burger tells the reader: here is one way to see purchasing power. It does not say: here is the final answer. It gives the mind a hook. Once the hook is in place, the heavier machinery can arrive: exchange rates, local wages, non-tradable services, production costs, inflation and income levels. The strange measure has done its job if the reader is now willing to follow the serious idea.

Other consumer-object indices try to borrow the same trick. A Starbucks latte can compare a familiar urban indulgence across places. An IKEA Billy bookcase can turn furniture into a price-level signal. These examples are charming because they are concrete. They are also noisy. A latte is not just milk, coffee and rent. It is brand positioning, city-center foot traffic, wage structure and the cultural willingness to pay too much for hot liquid in a paper cup. A bookcase is not just wood or particleboard; it is shipping, retail strategy and local tax. [note 2] [note 3]

The temptation is to keep adding objects until a proper basket appears: burgers, lattes, haircuts, taxis, hotel rooms, furniture, phones. In one sense, this is progress. A basket is generally more robust than a single product. In another sense, the charm begins to leak away. The more complete the basket becomes, the closer it moves toward official price statistics. The strange measure becomes less strange and more correct. That may be better for analysis, but worse for memory.

There is a tradeoff here that runs through the whole book. A single object is vivid but fragile. A full basket is robust but dull. The art is knowing which job you are asking the measure to do.

If the job is formal policy, do not use a cheeseburger as your main instrument. If the job is to make a person understand purchasing power in one clean motion, the cheeseburger has earned its ridiculous little crown.

That is the first lesson: a strange index is a translation device before it is a measurement device.

The best translation is not always the most complete. It is the one that gets enough of the meaning across before the listener walks away.

The Big Mac Index survives because it gives abstraction a bite mark. It takes the elegant but bloodless idea of purchasing power parity and gives it pickles. It also lowers the entry fee to argument. A reader cannot personally audit a national price basket before breakfast, but they can imagine the burger, ask whether McDonald's is expensive locally and disagree with the measure without being excluded from the conversation.

That is the best defense of popular measures. They are not perfect; they are invitations to better questions. The alternative is not a population calmly absorbing full methodology notes. The alternative is often no understanding at all.

The Big Mac creates several. Why does the same product cost different amounts in different countries? Why do exchange rates not equalize the experience of buying things? Why do local wages, rents and taxes matter? Why does a product's cultural status affect price? Why does a single object teach more quickly than a full basket?

Those are not childish questions. They are the right questions with ketchup on them.

The burger also reveals a hidden aesthetic rule of good indices: the unit should be slightly embarrassing. Not humiliating, just disarming. A measure that looks too official can intimidate the reader into passive acceptance. A measure that looks slightly absurd invites inspection.

The trouble begins when the embarrassment disappears. Once the Big Mac Index became famous, it acquired a small institutional dignity. It entered textbooks, investor notes and public argument. That is deserved, but it also changes the reader's posture. The joke becomes a brand. The brand becomes a reference. The reference begins to sound more authoritative than the method should allow.

This is how even good strange indicators age. They begin as explanations and become objects of belief. A measure that

once taught skepticism can become another thing people repeat without checking. The way to keep the Big Mac useful is to keep it a little ridiculous. The burger should never be allowed to put on a suit and pretend to be the International Monetary Fund.

Its commercial contamination matters. In formal measurement, contamination is often a flaw. In popular explanation, some mess can be a feature. A Big Mac price contains wages, rent, supply chains, brand strategy, taxes and consumer willingness to pay. Ordinary life is not pure. A measure that carries some of the mess may explain more than a measure that cleans itself into irrelevance.

This does not mean any messy object can become an index. A bad proxy is just clutter with ambition. The object must be comparable enough, repeated enough and understandable enough. It must be produced by a process that can be described. It must vary for reasons that are at least partly related to the thing being discussed. And it must fail in ways the reader can grasp.

That last condition is rare. The failure mode of the Big Mac Index is easy to explain: it is one product, sold by one company, affected by local conditions. The failure mode of many complex composite rankings is harder to explain because the formula contains so many choices that criticism itself becomes a profession. A simple proxy can be more honest because its weaknesses are visible.

The practical question is not whether a strange index is “serious” in the abstract. The question is what job it is doing.

If the job is conversation, the Big Mac is excellent. If the job is a first intuition about purchasing power, it is useful. If the job is formal currency valuation, it is insufficient. If the job is lunch, it is arguable.

The same measure can pass one exam and fail another. That does not make it fake. It makes it scoped. Public conversation hates scope. It wants verdicts.

The burger refuses, or should refuse, to provide one. It is at its best when it says: look here, then look deeper.

There is also a hidden lesson in the choice of a branded object. Economists normally prefer categories: food, housing, energy, services, durable goods. Categories are sensible because they reduce the weirdness of individual products. But readers live through individual products. They do not buy “tradable food item adjusted for local service inputs.” They buy lunch. A category may be more accurate, but a named object is easier to remember.

This is why the Big Mac Index belongs to popular economics rather than formal measurement. It is not trying to beat an official statistical basket. It is trying to translate the first step of the idea. If you begin with methodology, many readers leave before the idea has a pulse. If you begin with the burger, the reader already knows the object before learning the theory.

There are bad versions of this technique. A writer can use the object as decoration and never do the intellectual work. The burger becomes a gimmick, the lipstick becomes a joke, the underwear becomes clickbait. The measure appears in the headline but disappears in the analysis. This book cannot do that. The object must always pay rent. It must open a real question, not merely decorate one.

The Big Mac pays rent because it sits exactly where multiple economic forces meet. It is standardized but local. It is global but priced in local currency. It contains traded ingredients and non-tradable services. It is familiar but not universal. It is cheap enough to be mass-market and complex enough to contain labor, rent, logistics and tax. It is a small thing with a surprisingly large shadow.

That does not mean every Big Mac chart deserves trust. The data still has dates. The comparison still has base currencies. GDP-adjusted versions make different claims from raw versions. Exact contemporary rankings need fresh source checks; the durable lesson is the method, not a stale leaderboard. Every currency comparison also has a point of view. Every base country makes some experiences look normal and others look strange.

If a burger in one country looks cheap from the outside, it may not feel cheap inside the wage structure of that country. This is the bridge to Chapter 2. A currency comparison is not yet a life comparison. The Big Mac can show that exchange rates and local prices diverge, but it cannot tell us how much of a worker's day the burger consumes. For that, the unit must change.

The burger opens the door. Time walks through it.

There is a nice historical irony in using a mass-produced American sandwich to teach international price theory. The Big Mac is a symbol of standardization, yet the index exists because the standardized thing refuses to have a standardized economic meaning. Same brand, different burden. Same menu item, different currency story.

That is the kind of contradiction a good chapter can use. The product's entire commercial identity says "you know what this is." The price comparison says "not really." You may know the object, but you do not know what it means in each economy until you ask how it is priced, earned, taxed, supplied and perceived.

This is why the burger is stronger than a more neutral product might be. A generic basket of groceries would be more representative but less narratively sharp. A bus ticket might be locally important but less global. A luxury watch might travel but would speak to elites. The Big Mac sits awkwardly in the middle: common enough to be legible, standardized enough to compare, compromised enough to be interesting.

The burger does not need endless apology. The better rhythm is confidence followed by calibration. Show the burger doing its trick. Then show the trick's limits. The reader must first see why the measure became famous. Only then will the limitations feel like intelligence rather than pedantry.

This rhythm will repeat throughout the book. Let the measure work. Then inspect the mechanism.

The next point is why the Big Mac is not merely a funny example of PPP, but a gateway into comparative humility.

People often speak about currencies as if they have an obvious correct value. Strong currency, weak currency, cheap country, expensive country. These phrases sound clean until one asks: compared with what? A currency may look cheap against the dollar and still buy little for local workers. A country may look expensive to a tourist and ordinary to a resident with local income. A single exchange rate contains trade, capital flows, interest rates, politics, expectations and fear. It is too much for one sandwich.

But the sandwich can break the spell of nominal exchange rates. It says: do not look only at the money conversion. Look at what money buys. That single shift is large enough to justify the index's fame.

The burger also introduces the difference between analytical seriousness and communicative seriousness. An academic model can be analytically serious and communicatively inert. A burger can be analytically rough and communicatively serious. The Big Mac Index is good because it does both well enough for its purpose. It has published data, a repeatable idea and obvious caveats. Those three qualities will become part of the reader's informal test for every later indicator.

Does it have data?

Does it have method?

Does it admit limits?

Many indicators in the book will fail one of those tests. The burger mostly passes.

And once you accept that a burger can translate money, the next question is obvious.

Why stop at money?

If prices are supposed to tell us what things cost, perhaps they are using the wrong unit. A price in dollars, euros or pesos is only half a sentence. The other half is the life required to earn it.

Chapter 2 — Prices Become Human

A thousand-dollar phone is not the same object everywhere.

Physically, it may be close enough. Same glass, same camera bump, same small ritual of pretending you will not drop it. Economically, it can be several different creatures. In one country it is a mildly painful purchase. In another, it is a month of life. In a third, it is a luxury artifact that has escaped from the electronics section and wandered into social anthropology.

Converted prices hide this. Work-time comparisons reveal it.

The shift is simple: stop asking what the object costs in dollars and ask how long someone has to work to buy it. Suddenly a phone is not a phone. It is Tuesday, Wednesday and half of Thursday. A burger is not a burger. It is ten minutes in one city and more than an hour in another. The unit has changed from money to life.

That is why work-time indices feel more human than sticker prices. They translate economic conditions into the one budget nobody can print more of.

Reports such as UBS Prices and Earnings have used city-level price and wage comparisons to show how long workers in different places must work to afford the same goods. The method depends on gross or net hourly earnings, annual working hours, tax assumptions and the chosen product price. [note 4] Secondary summaries have applied similar logic to products such as iPhones. [note 5]

This sounds clean. It is not.

The first problem is the worker. Which worker? Average wage, median wage, skilled wage, service wage, formal-sector wage, net wage after tax, gross wage before deductions? A phone bought by a software engineer and the same phone bought by a delivery driver do not represent the same burden. One number will always flatten someone.

The second problem is the product. A phone model may

be globally recognizable, but prices can include taxes, carrier subsidies, import duties, warranty structures and promotional distortions. A cheap-looking headline can turn expensive once local financing and income are considered. An expensive-looking headline can become ordinary in a high-wage city.

The third problem is the city. City comparisons are seductive because cities feel concrete. Zurich versus Mumbai. New York versus Nairobi. Tokyo versus Mexico City. But cities are not countries, and a central-city wage is not a national life. A comparison can be excellent at showing urban purchasing power and useless as a statement about everyone living under the same flag.

So why use the measure?

Because the alternative is worse. Raw price comparisons are often stupid in a very polished way. A product that costs \$900 in two places has not imposed the same burden if one buyer earns \$40 an hour and another earns \$4. The dollar has done its arithmetic correctly while still failing to explain the situation.

Work-time fixes that failure. Not perfectly, but violently enough to be useful.

It also has a moral clarity that economists sometimes dislike because it is hard to hide from. Saying that a phone costs \$1,000 sounds like a consumer fact. Saying that it costs three weeks of ordinary labor sounds like a social fact. The object has not changed. The measurement has.

The same logic applies to smaller things. A burger can be translated into minutes. A cup of coffee can be translated into slices of a working day. A date night can be translated into the amount of labor required to perform the rituals of urban romance: food, drinks, transport, tickets, perhaps one desperate attempt to look relaxed while checking the bill. [note 6]

These measures are funny because they are intimate. They are serious for the same reason. A cost-of-living index is abstract until it tells you that a normal evening out consumes

more of your working week than you expected. People understand affordability when it steals time from them.

But work-time comparisons can also flatter the analyst. They look humane, therefore they feel truthful. That is dangerous. The method still needs discipline.

A good work-time comparison states its wage source. It says whether wages are gross or net. It names the product and price source. It avoids pretending that one occupation represents all workers. It does not update a global ranking once and then let the numbers rot for five years while the prose continues wearing them like fresh clothes.

The best version of the measure is not “Country X is rich and Country Y is poor.” That is not insight; that is a headline with a passport. The better version is: the same object enters different lives at different weights.

That idea changes how we read prices.

A latte in one city may be a casual daily indulgence. In another, it is a deliberate expense. A haircut may be cheap in dollar terms and expensive in time terms. A taxi ride may look inexpensive to a visitor and ruinous to a local worker. Rent is the extreme case: it is not merely money paid for shelter, but hours of life transferred to a landlord with suspiciously little ceremony.

The work-time lens also explains why people argue so bitterly about affordability. They are often measuring different realities. One person sees a nominal price. Another sees wages. Another sees rent. Another sees debt. Another sees social expectation. A single object becomes a little courtroom where the whole economy is on trial.

Strange indicators help when they reveal which reality is being discussed. They mislead when they smuggle one reality in as all of them.

The iPhone is a good example because it is absurdly modern and ancient at the same time. Modern because it is a globally distributed slab of electronics. Ancient because the underlying question is old: how much of my labor must I trade for this tool, status symbol, communications device and anx-

iety rectangle?

That question is more honest than asking whether the phone is “expensive.” Expensive compared to what? Other phones? Last year? Local wages? The buyer’s rent? The amount of time the buyer will spend staring at it instead of living a calmer life?

Work-time indices cannot answer all of that. They answer enough to make the abstraction human.

There is a reason time feels more honest than money.

Money is already an abstraction. It changes costume across borders and classes. One person’s “cheap” is another person’s postponed dental appointment. One household’s normal subscription stack is another household’s weekly food budget. Money pretends to be a common language, but in lived experience it is full of accents.

Time is not perfectly equal either. An hour of pleasant work and an hour of exhausting work are not the same. An hour for a parent juggling childcare is not the same as an hour for someone with no dependents. An hour in a safe job is not the same as an hour in precarious work. But time has a brute quality that money lacks. It reminds us that affordability is not only arithmetic. It is sacrifice.

This is why “hours to buy” comparisons can feel almost rude. They remove the soft cushion around price. A luxury object can no longer hide behind branding if the reader sees it as forty hours of work. A cheap meal can no longer look cheap if it consumes a meaningful slice of local wages. A phone can no longer be discussed as one global product if the work required to buy it changes dramatically by place.

The measure also exposes a strange feature of modern global culture. Advertising globalizes desire faster than wages globalize purchasing power. The same phone launch, the same sneaker, the same streaming show, the same lifestyle image reaches people whose local earning power differs wildly. The object becomes globally familiar before it becomes globally affordable. Work-time comparisons make that gap visible.

That is why the iPhone is such a strong example. It is not merely expensive. It is symbolically synchronized. People around the world know the product at roughly the same time. They see the same camera improvements, the same keynote language, the same smooth videos of aluminum being treated as if it had a spiritual life. But the labor required to buy the object is not synchronized. The marketing is global. The affordability is local.

This is one of the cleanest examples of what strange measures can do. They reveal a mismatch between the story a market tells and the life a buyer inhabits.

A normal price tag tells you what the seller asks.

A work-time price tells you what the buyer surrenders.

Those are not the same sentence.

There is also a political edge here, though not a policy tract. When people argue about whether life has become more expensive, they often throw around nominal prices, inflation rates and personal anecdotes. Work-time comparisons can cut through some of that. If a product takes fewer hours to buy than it did twenty years ago, something has improved in that narrow sense. If rent, care, education or medical costs take more hours, something has worsened in a more important sense. The method forces specificity.

It also prevents nostalgia from doing all the work. People are terrible at remembering prices. They remember the emotional price of being young and confuse it with the economic price of goods. Work-time can correct some of that by asking: compared with wages, what happened?

But the measure can be cruel if used carelessly. It may imply that everyone has access to the wage used in the calculation. It may flatten informal labor, unpaid care, unemployment or underemployment. It may compare an average wage to a product bought disproportionately by people above that wage. It may turn inequality into a single neat number and then congratulate itself for being humane.

The answer is not to abandon the measure. The answer is to write it with manners.

Say whose time. Say which wage. Say which price. Say what the comparison leaves out.

This is not academic fussiness. It is the difference between a useful translation and a misleading one.

Work-time also highlights why multiple strange indicators can complement one another. The Big Mac shows purchasing power through a product. The iPhone shows global aspiration through labor burden. A rent-to-burger comparison, if built carefully, might show how shelter distorts ordinary consumption. A date-night comparison might show social life under inflation. None is complete. Together, they create a better texture of affordability than a single headline price.

But the examples need progression, not hoarding. The argument does not need every affordability comparison in the world. It needs the conversion from price to life to land. Enough to show the principle. Not enough to become a municipal price database with jokes.

The strongest insight is simple: affordability is not what something costs in money. It is what it costs in the life available to the person buying it.

That sentence should remain underneath the whole chapter, even if the prose never states it so neatly.

Work-time also has a literary advantage: it produces instant inversion.

A city that sounds expensive may become less frightening once wages enter. A city that sounds cheap may become harsher once local earnings are considered. This inversion is useful because it breaks the tourist's view of the economy. Visitors often judge places through converted prices. Residents live through local incomes. The visitor says, "Coffee is cheap here." The resident may hear, "Your hour is cheap here." The same cup has acquired a different moral temperature.

The tone matters. Work-time comparisons can easily become poverty tourism if written badly. The purpose is not to gawk at how many hours someone else must work for a phone. The purpose is to show that prices only become so-

cially meaningful when placed beside income, time and necessity.

The hierarchy of goods matters too. A phone is interesting because it is globally desirable and widely understood, but it is not survival. Food, rent, electricity and transport sit closer to necessity. A date-night index is amusing because it mixes economics with social ritual. Rent-to-burger comparisons are stranger because they reveal how housing can distort ordinary consumption. The point is how the unit changes the feeling: dollars create distance, hours create pressure, necessities create urgency, social rituals create recognition.

There is an especially useful contrast between a phone and a burger. The burger is repeated, perishable and local. The phone is durable, global and aspirational. The burger teaches purchasing power through everyday consumption. The phone teaches inequality through synchronized desire. Both are consumer objects, but they do not perform the same conceptual job.

This helps prevent the book from becoming “fun facts about prices.” Measurement is an act of framing. A price framed as currency produces one response. The same price framed as work produces another. A good index is not merely a calculator. It is a choice about which human experience should be foregrounded.

Later, with AI benchmarks, the same lesson will return in a colder form. A model score framed as percentage accuracy produces one response. The same capability framed as hours saved, failure rate or unsupported autonomy risk produces another. The unit shapes the story. Chapter 2 gives the reader that insight in human terms before the technical examples arrive.

It also gives the book its first emotional escalation. Chapter 1 is clever. Chapter 2 is personal. The Big Mac asks whether currencies make sense. Work-time asks what the economy takes from the person inside it. A city can look affordable under one wage and brutal under another. The measure can teach inequality even when it is not designed for that

purpose.

One elegant way to handle this in prose is to present work-time as a question rather than a verdict. “How many hours does this cost?” naturally invites “whose hours?” That second question opens the door to wage distribution, employment type and class without requiring a technical detour.

This is what a useful strange measure does: it creates the next honest question.

City rankings are tempting and usually distracting. Rankings are seductive but fragile. They age quickly, depend on data choices and distract from the conceptual point. A few examples can illustrate the spread; the durable idea is the method.

A good work-time measure changes ordinary perception. The next time they see an expensive object, they may ask not “what does it cost?” but “how much life does this demand from different people?” That is a useful habit. It does not solve economics. It makes economics harder to hide.

There is another useful pressure point: work-time comparisons make inequality visible without requiring the reader to start with inequality theory.

Take a product that appears universal. A phone, a meal, a pair of shoes, a train ticket. The object enters a city wearing the same brand name. Then work-time removes the disguise. For one worker, the object is a casual upgrade. For another, it is a planned sacrifice. For a third, it is simply outside reach. The product is global; the burden is stratified.

This also explains why some consumer debates become emotionally explosive. People are not only arguing about the product. They are arguing about whether the economy recognizes their time. A person told that an object is affordable may hear that their labor is being dismissed. A person told that a city is cheap may hear that local wages have vanished from the conversation.

The measure makes that insult harder to miss.

This is why work-time belongs early in the book. It gives the reader a humane distrust of averages. It teaches that the

same price can produce different realities before the book moves into more theatrical examples. Without that grounding, lipstick and underwear could become mere jokes. With it, they become stories about stress, substitution and hidden cuts.

Time is not the final truth either. Time is powerful because it is intuitive, but it is not everything. A household with savings experiences a purchase differently from one without. A person with debt, dependents or insecure income experiences the same hour differently from a person with stability. Work-time improves the price story; it does not complete it.

This is a useful lesson for the whole book: a better proxy is still a proxy.

The practical result is suspicion of price tags. A price is never just a price. It is a claim against someone's time, income, obligations and alternatives. The number on the shelf is the seller's language. The hours behind it are the buyer's translation.

They also shift the book's center of gravity. Chapter 1 showed that an object can translate currency. Chapter 2 shows that an object can translate life. Once a reader sees that, a stranger class of indicators becomes possible. If a burger can explain exchange rates and a phone can explain wages, perhaps lipstick can explain recession anxiety.

This is where the subject becomes more entertaining and less trustworthy.

Prices made human are still prices. They retain a mathematical spine. Folklore indicators often have something softer: a story that sounds so right people forget to ask whether it is true.

That is how the underwear drawer enters economics.

Chapter 3 — The Folklore of Recession

There is a wonderful kind of economic indicator that sounds plausible for about eleven seconds.

The Lipstick Index is one of them. In hard times, the story goes, people stop buying expensive luxuries but still allow themselves small ones. Not the handbag, but the lipstick. Not the new wardrobe, but the little tube of color that says civilization has not entirely collapsed. It is neat, vivid and just faintly glamorous. It also has the excellent commercial advantage of making recession sound like a department-store problem.

The Men's Underwear Index is less glamorous, which may be why people remember it even more. The idea is that men replace underwear less often when the economy weakens because nobody sees it and the purchase can be postponed. It is popularly attributed to Alan Greenspan. It has the useful smell of secret central-bank eccentricity: while ordinary people watch GDP, the monetary priesthood studies boxer shorts.

Then come hemlines, skyscrapers and other creatures from the same forest. Skirts allegedly rise and fall with markets. Record-breaking towers supposedly appear at the edge of downturns. Each story has a shape the mind enjoys: private behavior reveals public stress; vanity exposes macroeconomics; architecture confesses hubris before finance does.

The problem is that a good story can outrun its evidence.

The Lipstick Index is best treated as a media-friendly hypothesis, not a dependable recession tool. Leonard Lauder's early-2000s observation helped popularize the idea, and later writers have revived it whenever cosmetics sales, consumer anxiety and recession talk happen to stand near one another. But subsequent evidence is mixed, and some explanations shift from lipstick to nail polish, mascara or other affordable luxuries depending on the moment. A theory that keeps changing cosmetics bags may still contain insight about consumer psychology, but it is not a stable macroeconomic instrument. [note 7] [note 8]

The Men's Underwear Index has a similar charm problem.

There is limited empirical discussion of whether men's underwear sales predict economic conditions, and the Greenspan attribution is repeated more confidently than it is usually sourced. The underlying behavior is plausible: hidden replacement purchases can be delayed. But plausible is not the same as measured, and memorable is not the same as true. [note 9]

This distinction matters because folklore indicators do not feel weak. They feel better than ordinary statistics. They feel like hidden knowledge. That is their commercial power. An inflation report is public. An underwear index feels like a trick the serious people forgot to hide.

Economic folklore has three ingredients:

First, it is concrete. Lipstick, underwear, skirts and skyscrapers all make better mental pictures than consumer sentiment surveys.

Second, it is psychologically satisfying. It tells us that people reveal the economy through small involuntary acts. They may deny anxiety, but their shopping carts confess.

Third, it is just hard enough to disprove in casual conversation. Sales can be sliced by category, market, income group, brand, season and cultural moment. One downturn may appear to support the story; another may embarrass it. By the time the evidence has been made precise, the anecdote has already left the room and got itself invited onto a podcast.

The Skyscraper Index belongs to the grander wing of the same building. It suggests a connection between the construction of record-breaking skyscrapers and economic downturns. The symbolic appeal is irresistible. Societies build upward at the moment confidence becomes delusional. Cheap money, credit expansion and civic vanity combine to produce a tower that opens just in time for everyone to discover the lifts are not the only things going down.

This is marvellous prose and dangerous analysis.

There may be useful historical and cycle observations around construction booms, credit and property markets. Tall buildings do not grow out of sober pessimism. But the

stronger claim — that skyscrapers reliably predict recessions — is contested. Research and critiques have treated the relationship as more correlation story than forecasting law. The image is powerful; the measure is fragile. [note 10] [note 11]

This is not an exercise in sneering at folklore. Folklore often preserves something real. People do buy small comforts under stress. Hidden purchases can be postponed. Construction booms can say something about confidence and credit. Clothes, cosmetics and towers all belong to the economy because everything does. The issue is not whether the stories are stupid. The issue is whether we make them carry more weight than they can bear.

A folklore indicator can be useful as a question:

What are people quietly cutting?

Which small luxuries survive anxiety?

Which visible monuments appear near the top of a cycle?

Where does confidence become performance?

Those are good questions. They become bad answers when turned into prediction machines.

The safest prose around these indicators uses verbs like “suggest,” “reflect,” “hint,” “popularly attributed,” “often repeated,” “contested” and “weakly evidenced.” This is not cowardice. It is respect for the difference between a signal and a story about a signal.

The reader may reasonably ask why we keep these indicators at all. If they are weak, why not throw them out and return to respectable data?

Because they teach the central lesson better than respectable data can. They show how badly we want the economy to reveal itself through ordinary things. They show that the appetite for strange measures is not only analytical. It is emotional. People want an early warning. They want a private clue. They want to believe that a tube of lipstick or a packet of underwear will whisper what official statistics say too late.

That desire is not irrational. Official statistics can lag. Formal models can miss lived experience. But folklore is not au-

tomatically better because it is vivid. A myth can travel faster than a spreadsheet and still arrive wrong.

By now the pattern is becoming clearer.

Folklore indicators also reveal something uncomfortable about expertise.

Experts often dislike them because they are messy, weakly evidenced or media-friendly in precisely the wrong way. Fair enough. But the public often likes them because they feel closer to life than the expert vocabulary does. The expert says “consumer confidence weakened.” The folklore indicator says “people stopped buying new underwear.” The expert is more rigorous. The underwear is more memorable. The gap between the two is where popular economics lives.

The responsible writer cannot simply side with one camp. If we side only with expertise, the result becomes a scolding lecture about why fun things are wrong. If we side only with folklore, the book becomes a circus of charming falsehoods. The useful position is less comfortable: take the folklore seriously as culture, but not necessarily as measurement.

The Lipstick Index, for example, may be weak as a recession predictor and still strong as a story about affordable luxury. It captures a real behavioral possibility: when large pleasures become unaffordable, small pleasures may become more meaningful. That does not mean lipstick sales forecast GDP. It means lipstick can help readers think about substitution, mood, identity and consumption under stress.

The Men’s Underwear Index may be weak as formal economics and still reveal an important class of spending: private, postponable, replacement-based purchases. Many households do not cut spending evenly. They delay what can be delayed. They keep visible obligations and quietly stretch invisible ones. Underwear is funny because it is hidden. The economic idea behind it is not funny at all.

Skyscrapers may fail as prediction machines and still say something about credit, optimism and monument-building. A record tower requires capital, confidence, political permission, design ambition and a long enough boom for everyone

involved to believe the sky is available. The problem is not that skyscrapers have no economic meaning. The problem is the leap from meaning to reliable forecast.

This is a recurring pattern in bad index thinking: a measure begins as metaphor, then sneaks into prediction.

Metaphor is allowed to illuminate one aspect of reality. Prediction must answer to evidence. Confusing the two is how cleverness becomes misinformation.

The Hemline Index makes this especially clear. The idea that skirt lengths move with economic cycles has the delightful quality of sounding like something a columnist invented after lunch and then refused to abandon. It may reveal attitudes, fashion cycles, media habits or the human desire to detect hidden patterns. But as a core economic indicator, it is too fragile for this book. It belongs in the optional drawer, clearly labeled folklore, not in the engine room of the argument.

This does not make it useless. Folklore examples give the book texture. They show the reader how attractive weak measures can be. They provide contrast with stronger examples later. A chapter that included only robust measures would teach less about myth. A chapter that included only myths would become frivolous. The balance matters.

Entertainment is the bait; skepticism is the hook. The reader can enjoy the absurdity of recession indicators hiding in cosmetics counters and underwear drawers, then gradually feel the seriousness of the evidence problem.

The danger is tone. Too much mockery and the result becomes smug. Too much indulgence and it becomes credulous. The voice can stay amused without becoming superior. The invitation is better judgment, not punishment for liking a good story.

The structure can do this naturally: open with the irresistible claims, show why they make psychological sense, then ask what the evidence can bear. Finally, explain why the stories survive anyway.

The survival question is crucial. Weak indicators do not

survive only because people are foolish. They survive because they satisfy needs that formal indicators do not. They are portable. They are visual. They let people feel ahead of the official narrative. They turn macroeconomics into gossip, and gossip is one of humanity's most successful information technologies.

That does not make the gossip true. It explains why it spreads.

The useful instinct is simple: when an index sounds too perfect as a story, slow down. The more delicious the anecdote, the more carefully it should be handled.

This will prepare them for Chapter 4, where the measures become less charming but more operational. A shipbroker's route assessment will never beat lipstick for cultural sparkle. It may, however, have a better claim on reality.

That is the tradeoff the book keeps returning to. The world gives us signals. Some are vivid. Some are grounded. The rare ones are both.

There is also a gender and culture trap here.

Lipstick, underwear and hemlines invite lazy writing. It would be easy to turn them into jokes about women shopping, men being oblivious and markets reading our wardrobes. That would be cheap and probably boring. The better approach is to treat the objects as consumer-behavior stories shaped by culture, marketing and visibility. Lipstick is not interesting because women are frivolous. It is interesting because small luxuries can survive stress. Underwear is not interesting because men are ridiculous, although evidence suggests they often help the argument. It is interesting because hidden replacement purchases can be deferred.

The joke should punch at weak claims, not at the consumers inside them.

This matters commercially too. A popular nonfiction book can be sharp without being lazy. Readers enjoy wit, but they can feel when the joke has no intellectual spine. The best humor here comes from proportion: the immense seriousness of recession anxiety filtered through the tiny absurdity of lip-

stick and underwear. The objects are funny because they are too small for the burden people place on them.

The line between folklore and misogyny, folklore and class snobbery, folklore and “people are stupid” must stay visible. The book’s stance is not that ordinary consumer behavior is foolish. Ordinary behavior is often where macro pressure becomes legible. The foolishness is in overclaiming.

That gives the argument a generous skepticism. It can say: there is a real human impulse here, but not enough measurement discipline. That is more interesting than saying: ha, silly index.

Recession folklore also needs a distinction between three related but different claims.

The first is behavioral: people change small purchases under economic stress. This is plausible and often observable in specific categories.

The second is diagnostic: those changed purchases can reveal something about consumer mood or constraint. This may be true in context.

The third is predictive: those purchases reliably forecast recessions. This is the dangerous leap.

Many weak indicators survive by sliding between these claims. A writer begins with behavior, moves to diagnosis and ends with prediction. The reader feels continuity because the object remains the same. But the evidentiary burden has changed completely.

This is a useful trick to expose. It applies beyond lipstick. A rising skyscraper can indicate credit optimism without predicting a crash. Underwear sales can reflect household caution without functioning as a macro dashboard. Hemlines can be an amusing cultural anecdote without becoming a market signal. The same story can be valid at one level and overclaimed at another.

The disciplined question is: what strength of claim is being made?

Is the writer saying “this can happen”?

“This happened once”?

“This often happens”?

“This predicts what happens next”?

Those are different animals. They should not share a cage.

This claim-strength discipline is one of the quiet skills the book can give readers. It is also a way to make skeptical prose more precise. Instead of saying “the Lipstick Index is nonsense,” the book can say: the affordable-luxury behavior is plausible; the recession-predictor claim is contested. That is sharper, fairer and more useful.

There is one more reason folklore indicators matter: they show how economic stories compete for attention.

A serious indicator must often explain itself before it can be understood. A folklore indicator arrives pre-explained. Lipstick goes up because people want small luxuries. Underwear sales fall because men delay invisible purchases. Skyscrapers rise because boom-time arrogance wants a monument. Hemlines move because culture and money are supposedly dancing together. The causal story is built into the object.

That built-in story is the advantage. It is also the evidence hazard.

A good causal story can make weak data feel stronger than it is. The reader hears the explanation and supplies belief. This is a form of narrative credit. The indicator borrows trust from plausibility. That is why contested indicators need careful handling in a popular book. If the prose is too smooth, it will accidentally upgrade them.

One way to prevent that is to make uncertainty part of the entertainment. The prose can invite the reader to enjoy the absurdity while watching the evidence wobble. The humor should come from the gap between confidence and support. Lipstick may tell us something about consumer psychology; it probably should not be allowed to forecast industrial production. Underwear may reveal deferred replacement spending; it does not deserve a marble plinth in the temple of macroeconomics.

The useful lesson is that not all weak indicators are equally

useless. Some are useless as forecasts but useful as cultural diagnostics. The Lipstick Index may be poor at predicting recessions and still excellent at revealing how people talk about recession. The Skyscraper Index may be weak as a timing tool and still useful as a story about credit, ambition and late-cycle optimism. The Men's Underwear Index may be thin as data and still memorable as a lesson in hidden consumption.

This distinction lets the book keep the fun without sacrificing standards. It also gives the book a recurring move: rescue the interesting idea from the exaggerated claim.

That move is better than debunking for its own sake. Debunking can become tedious because it lives in the negative. The goal is not to stomp around saying "false" at every charming story. The goal is to ask what the story can responsibly do.

Can lipstick introduce the idea of affordable luxuries? Yes.

Can underwear introduce postponable hidden consumption? Yes.

Can skyscrapers introduce boom-time capital intensity and confidence? Yes.

Can any of them carry the claim "this predicts recessions reliably"? Not without much stronger evidence than public retellings usually provide.

This is the measure/myth boundary in action.

This also earns the title. Burgers may open the book, but lipstick and underwear give it its skeptical bite. Without them, the book would be a book about clever proxies. With them, it becomes a book about why clever proxies so easily become nonsense.

A good strange index has an object, a method and a limit.

A folklore index has an object and a story, then hopes nobody asks too many questions about the method.

The next family of measures is different. It is less intimate, less funny and much harder to fit in a shopping bag. It does not come from cosmetics counters or underwear drawers. It comes from oceans, ports, containers, coal, grain, iron ore and electricity.

If lipstick is the economy as anecdote, shipping is the economy with its sleeves rolled up.

Chapter 4 — The Economy Under the Economy

Most consumer indicators begin at the shelf.

Shipping indicators begin before the shelf exists.

Before a shirt becomes fashion, cotton moves. Before a phone becomes desire, components move. Before a loaf of bread becomes breakfast, grain moves. The visible economy is full of packaging, brands, prices and moods. Underneath it sits a less photogenic economy of vessels, containers, ports, fuel, routes, capacity, warehouses, energy and waiting.

It is not cute. That is one reason to trust it slightly more.

The Baltic Dry Index is one of the best-known signals from this deeper layer. It tracks the cost of moving major dry bulk commodities by sea: things like iron ore, coal and grain. Unlike lipstick or underwear, it is not a consumer anecdote masquerading as a secret. It is a market benchmark rooted in physical shipping. The Baltic Exchange describes its bulk freight indices as based on assessments of route fixtures or price points, taking into account previous transactions, market conditions, supply and demand, with vetted shipbrokers submitting route assessments. [note 12]

That does not make it magic. It makes it grounded.

Dry bulk shipping is tied to the movement of raw materials. If demand for iron ore, coal or grain transport changes, freight rates can respond. If vessel supply is tight, rates can jump. If trade slows, rates can weaken. The signal is closer to the physical pulse of production than a story about lipstick sales. Ships are expensive, slow-moving assets. You do not conjure new fleets overnight because a spreadsheet feels optimistic.

But even grounded measures have limits.

The Baltic Dry Index is not “the world economy.” It is a slice of the freight market. It reflects dry bulk shipping conditions, not services, software, consumer confidence, local housing or whether anyone can afford a dentist. It can be distorted by vessel supply, route-specific conditions, commodity cycles and shocks that are real but narrow. A grounded signal is still a signal from somewhere.

That “somewhere” is the point. A good index tells you what room it is standing in.

Container freight indices stand in a different room. A container is not dry bulk. It is the metal box that made modern retail feel inevitable. Container rates describe the movement of manufactured goods and components through trade lanes. Freightos Baltic Index data, for example, concerns 40-foot container spot freight rates, using aggregated market data and calculation methods distinct from dry bulk shipping. [note 13]

People often lump shipping indicators together because ships are involved, which is like lumping restaurants and hospitals together because both contain sinks. Dry bulk and container freight answer different questions. Dry bulk asks about raw materials and vessel classes. Container freight asks about standardized boxes, trade lanes and the flow of goods closer to consumer supply chains.

The difference matters because one of the recurring sins of index reading is category laziness. A measure becomes famous, then gradually expands in casual speech until it appears to measure everything. “Shipping is up” becomes “the economy is strong.” “Containers are expensive” becomes “inflation is coming.” Sometimes those statements may contain something useful. Often they are missing enough detail to be dangerous.

The same is true of electricity demand. Electricity is a beautiful signal because it is hard to fake at scale. Factories, homes, data centers, offices, cold storage and railways all need power. Changes in demand can say something about ac-

tivity, weather, industrial production or crisis. Energy-data organizations and operational maps can make that activity visible. [note 14] [note 15]

But electricity is not pure GDP in cable form. Weather matters. Efficiency matters. Industry mix matters. Solar generation and grid structure matter. A heatwave can make electricity demand look economically lively because everyone is trying not to melt. A recession can coincide with technological or structural shifts that muddy the picture. Again: useful signal, limited room.

Restaurant reservations are even more delicate. They can reveal consumer behavior quickly. They can also reflect seasonality, local rules, tourism, platform coverage and whether people have decided that paying restaurant wine markups is an act of financial self-harm. As a nowcasting clue, reservations can be interesting. As a grand index of economic truth, they should be kept on a shorter leash.

The physical-economy chapter is important because it restores seriousness after the folklore chapter. It shows that not all strange indicators are equally flimsy. Some are strange only because ordinary readers do not usually stare at freight assessments. A shipping index may sound obscure, but it can be methodologically clearer than a recession theory built out of lipstick.

There is an irony here. The more grounded indicator is often less culturally viral. Everybody can imagine lipstick. Fewer people can imagine a Capesize vessel market without their eyes politely leaving the conversation. The economy rewards vividness in attention, but not always in truth.

This creates a communication problem. The best indicators may need translation before they can reach readers. The Baltic Dry Index tells a stronger story once the reader understands that it is not an abstract financial number floating above the world, but a price connected to moving heavy things across water. Container freight becomes more vivid once the reader sees that the modern shop shelf is a sequence of ports and boxes pretending to be convenience.

The strange measure must make the invisible visible.

That does not mean it must be cute. It means it must be legible.

The physical economy also teaches humility. A freight rate can rise because demand is strong, because vessel supply is constrained, because a route is disrupted, because fuel costs shift, because ports jam, because geopolitics interferes, because the world has once again discovered that “globalization” is a comforting word for a machine made of many breakable parts.

This is why shipping measures are excellent antidotes to lazy abstraction.

Modern consumers are trained to experience logistics as absence. A product is available or it is not. Delivery is tomorrow or late. Shipping is free, which is one of the great retail lies of the age; it means only that the cost has been buried somewhere more convenient. The physical system disappears when it works. It becomes visible only when it fails.

Freight indices make the disappeared system visible before failure reaches the shelf.

They also remind us that not all economic signals are about consumers. A consumer-led view of the economy is natural because most people experience markets as buyers. But before buying comes movement. Raw materials must move. Components must move. Finished goods must move. Energy must flow. Containers must be in the right place, which they frequently are not because the world is inconveniently shaped and human demand lacks symmetry.

This makes physical indicators feel less democratic than burgers but more structural. Not everyone can check a shipbroker assessment. Not everyone follows container trade lanes. But the measures point to systems that consumer anecdotes can miss.

A recession may be felt in lipstick sales after households change behavior. A shipping rate may move when industrial demand or logistical stress changes earlier in the chain.

That does not make shipping predictive in a magical sense. It means the signal sits in a different part of the sequence.

Sequence matters. Strange indicators are often discussed as if they compete in one general talent show: which one predicts the economy best? That is the wrong question. A better question is: where in the system does this measure sit?

The Big Mac sits at retail price and currency comparison. Work-time sits between wages and consumption. Lipstick sits in consumer mood and substitution. Shipping sits in physical trade. Electricity sits in production, weather and infrastructure. AI benchmarks sit in technical evaluation and reputation markets.

Each is a window. A window is useful if you know which wall it is in.

The Baltic Dry Index also raises a useful issue about expert assessment. People sometimes imagine that a real index must be purely transaction-based, as if every useful number simply falls out of a machine. In many markets, especially complex physical markets, expert assessment remains part of benchmark production. That does not make the benchmark fake. It means governance, methodology and contributor quality matter. The Baltic Exchange's description of vetted shipbrokers and representative routes is precisely the kind of institutional machinery a serious reader should look for. [note 12]

This is another contrast with folklore. The Men's Underwear Index may have a clever behavioral story, but the method is often vague in public retellings. A freight benchmark may be less amusing, but it can usually point to a methodology document. Boredom is not proof of quality, but a transparent method is a good sign.

Container indices add the modern supply-chain lesson. The container is a strange object because it is boring by design. Its genius is standardization. It has no romance. It is a metal rectangle that quietly reorganized global trade. As an index object, it is almost the opposite of lipstick. Lipstick is psychologically vivid. The container is operationally deci-

sive.

The reader does not need to become a shipping analyst. They need to understand that the economy contains layers. Some layers speak through prices on shelves. Some through wages. Some through consumer habits. Some through freight and energy. A strange index is useful when it helps the reader locate the layer.

There is also a humility about physical systems that modern digital culture badly needs. Digital products often make scale feel frictionless. A file copies. A service launches. A chart updates. But physical markets remain stubborn. Ships move slowly. Ports jam. Energy grids strain. Droughts happen. Wars close routes. A container in the wrong place is not corrected with a keyboard shortcut.

Physical-economy indicators restore the cost of reality.

That is why the physical layer earns its place. Without it, the book risks feeling like a tour of consumer metaphors. With it, the reader sees that strange indicators can be industrial, operational and institutionally serious. The measure/myth tension becomes richer. Some measures are not myths because they are strange. Some are strange because most people are separated from the systems they measure.

The physical layer also complicates the reader's trust instinct.

It is tempting to trust a shipping index simply because it sounds hard. Ships are large. Ports are industrial. Freight has the smell of reality. Compared with lipstick, it feels like a grown-up has entered the room wearing steel-toed boots. That instinct is partly fair. Physical-market indicators often have stronger grounding than consumer folklore. But hardness is not the same as completeness.

The Baltic Dry Index can rise for reasons that do not mean broad prosperity. Vessel supply can tighten. Weather and geopolitics can disrupt routes. Commodity-specific demand can shift. Financial speculation and expectations can influence behavior around real capacity. The index is closer to physical reality than a department-store anecdote, but it still

needs interpretation.

This is a useful moment to introduce the difference between signal and explanation.

A signal says: something changed.

An explanation says: here is why.

Many bad readings of indices confuse the two. The number moves, and the observer immediately attaches a story. Freight up means demand boom. Freight down means recession. Reservations down means consumer weakness. Electricity up means industry recovery. Sometimes the story is right. Sometimes it is merely available.

The discipline is to let the signal provoke investigation rather than replace it.

This is particularly important because grounded indicators often attract people who want certainty. A consumer folklore story is obviously flimsy enough to invite caution. A shipping benchmark looks serious enough to seduce a different kind of overconfidence. The analyst may think: this is real market data, therefore my interpretation is real. But data quality and interpretation quality are separate things.

The right posture is respect without worship: shipping and energy signals deserve more attention than lipstick, but they still require scope, context and humility.

That is a more mature lesson than “some weird indicators are good and some are bad.” The better lesson is: some indicators are closer to real flows, but no indicator is free from context.

This will matter again in Chapter 7. AI benchmarks often look rigorous because they have datasets, scores and leaderboards. Like shipping indices, they are methodologically more concrete than folklore. Like shipping indices, they can still be overgeneralized. A benchmark can be real and misread. A freight index can be real and misread. Reality does not protect a measure from bad storytelling.

Chapter 4 gives the reader one of the book’s strongest spatial shifts. The first three chapters happen near the body: mouth, hand, pocket, wardrobe. Chapter 4 moves outward

to ports, oceans, grids and dining rooms seen as aggregate behavior. The reader can feel the scale change. The economy is no longer the thing bought. It is the route, capacity and infrastructure that make buying possible.

That shift helps the book avoid triviality. Without the physical economy, a skeptical reader might think strange indicators are mostly media inventions. Freight and energy show that the index habit also lives in serious markets, where benchmark construction affects contracts, expectations and operational decisions.

This is also a good place to introduce the idea that some measures are used by insiders before they become explainable to outsiders. The Big Mac Index is built for public understanding. Freight benchmarks are built for market participants. When public writers borrow them, they must translate in the opposite direction: from expert signal into ordinary explanation.

That translation has its own risk. The writer may oversimplify to make the signal dramatic. They may call BDI a recession predictor because “dry bulk freight route assessment under specific supply-demand conditions” lacks fireworks. The pressure to make expert indicators viral can turn serious measures into folklore. A grounded index can be degraded by bad public use.

This is a subtle but valuable point: myths are not only born from silly objects. They can be made from serious numbers by careless storytelling.

This physical layer connects backward and forward. Backward, it contrasts with folklore. Forward, it prepares for benchmark governance and AI evaluation. The reader learns that evidence quality is not determined by how funny the object sounds. It is determined by method, source, scope and use.

One more theme belongs here: latency.

Different indicators speak at different speeds. Official statistics often arrive after the fact because serious measurement takes time. Consumer behavior can shift quickly but

noisily. Freight rates can respond rapidly to market pressure but require expert interpretation. Electricity demand may move minute by minute but needs weather and structural context. Restaurant reservations can appear almost instantly but are limited by platform coverage and behavior.

A strange index is partly a bet on timing.

Does it tell us something before official data does?

Does it tell us something official data cannot?

Or does it merely tell us the same thing with better packaging?

This timing question helps separate useful alternative indicators from decorative ones. A shipping rate may be valuable because it sits upstream of retail experience. A reservation index may be valuable because it changes quickly. A lipstick story may be less valuable if it is noticed only after recession anxiety is already everywhere. A work-time comparison may not be early, but it reframes affordability in a way official inflation does not.

“Early warning” language needs restraint. Many indicators are bad at prediction. But timing remains important because people reach for strange measures when official language feels late, incomplete or too abstract. They want the world translated sooner.

That desire is understandable. It is also where many myths enter.

One number cannot explain all that. It can point.

This is what separates useful indicators from toy indicators. The useful one invites you to ask the next question. What cargo? Which route? Which vessel class? Which time horizon? Which market structure? What else moved at the same time?

The toy indicator invites you to repeat a clever sentence.

The economy under the economy is full of signals that deserve attention precisely because they are not designed for dinner-party entertainment. They are working measures. They live in industries where money changes hands and operational decisions depend on reference points. That does

not make them infallible. It makes their failure modes more specific.

The argument now moves from the consumer surface to the operational substrate. But there is another category of strange measure that does something different again. It does not translate money, time, mood or freight. It translates risk.

The air outside your window contains measurements most people cannot feel directly until the damage is already being done. Parts per million, micrograms per cubic meter, particulate matter, exposure windows: accurate language that does not easily enter the imagination.

So someone turns the air into cigarettes.

This is either brilliant or reckless.

Usually both.

Chapter 5 — Air You Can Count

Air is difficult to measure emotionally.

You can smell smoke. You can see a brown sky. You can cough. But much of the danger in polluted air is statistical, cumulative and invisible. Fine particulate matter does not arrive with a small invoice explaining its long-term health implications. It enters the body without ceremony. The measurement unit — micrograms per cubic meter — is precise and almost completely useless to ordinary intuition.

So the mind reaches for a translation.

One of the most striking translations is the cigarette-equivalent pollution measure. Instead of telling people that fine particulate matter is at a given concentration, it says roughly: breathing this air is comparable, in risk terms, to smoking some number of cigarettes. Berkeley Earth helped popularize a method for relating PM_{2.5} exposure to cigarette-equivalent harm. [note 16]

The attraction is immediate. Cigarettes are culturally loaded, familiar and bodily. People may not know what a

PM2.5 reading means, but they know a cigarette is not a vitamin. The comparison turns invisible risk into an object with ash on it.

This is exactly why it works.

It is also exactly why it can mislead.

A cigarette-equivalent is not a literal chemical equivalence. Polluted air and cigarette smoke are not the same exposure. The analogy is a risk translation, not a statement that the body has received an identical dose of the same substances. The number compresses complex epidemiology, exposure assumptions and health risk into a familiar unit. That compression is powerful because it is memorable. It is dangerous because memorable health analogies tend to harden into literal claims as they travel.

Chapter 5 is not an environmental-health guide. It is a chapter about human-scale measurement. Air pollution matters here because it gives the book one of its clearest examples of abstraction becoming readable. The story is not “air is bad.” The story is: a strange unit can make a hidden risk understandable, and the same act of translation can oversimplify the risk.

That distinction must stay bright.

There are several measures in this area, and they are easy to confuse. PM2.5 concentration is a measurement of fine particulate matter in the air. AQI is an index designed to communicate air quality and health concern. Annual exposure is not the same thing as a short-term spike. A daily average is not the same thing as the worst hour. Different countries can use different AQI methods, and OpenAQ explicitly notes that air quality indices are not standardized globally. [note 17] AirNow explains the U.S. AQI as a public communication index with categories and health concern levels. [note 18]

When these measures are mixed casually, the prose becomes dangerous. A writer may take a PM2.5 number, an AQI category, a daily reading and a long-term exposure claim, toss them into the same paragraph and produce something that sounds urgent but no longer means anything precise.

That is not popular explanation. That is fog with footnotes.

The cigarette analogy demonstrates the central bargain of this book:

Make the measure vivid enough to matter.

Keep it honest enough not to become nonsense.

The best version of the cigarette-equivalent idea says: this is a way to visualize risk. It does not say: this city literally smoked X cigarettes for you today. It does not turn one day's air into a lifetime diagnosis. It does not pretend that all bodies, locations, durations and pollutants collapse neatly into one tidy stick of tobacco.

There is a family resemblance here with carbon equivalents. A tonne of CO₂ is abstract. A flight is familiar. Carbon calculators often translate emissions into flight distances, car miles, household energy or other everyday comparisons. ICAO provides a carbon emissions calculator methodology for flights, and explainers such as Our World in Data discuss transport emissions in CO₂-equivalent terms, including assumptions beyond carbon dioxide alone. [note 19]

Again, useful and limited.

Carbon equivalents help because the climate system is too large for ordinary perception. Nobody feels one kilogram of carbon dioxide leaving an aircraft seat. But people understand a flight. They understand a car journey. They understand a household. The translation gives scale a shape.

Equivalence measures work best when the thing being measured is both important and emotionally unavailable.

Air pollution is like that. Climate emissions are like that. Radiation comparisons, calorie comparisons, water-footprint comparisons and energy equivalents all live in the same neighborhood. The underlying variables may be measurable, but they do not enter everyday intuition naturally. Nobody evolved to feel a microgram per cubic meter. Nobody wakes up with a native sense of atmospheric carbon concentration. The body understands smoke, heat, fatigue and smell; it does not understand a data table.

So the translator reaches for a familiar harm.

This move is powerful because it bypasses explanation. It says: you already know this kind of bad. Now use that feeling to understand this other kind of bad.

The danger is that feelings do not preserve methodology.

A cigarette is not just a unit of risk. It is an object loaded with moral, cultural and personal meaning. For some readers it suggests addiction, cancer, family illness, smell, carelessness or old glamour curdled into public-health warning. When a pollution index says “cigarettes,” it borrows all of that. The borrowing gives force. It also creates an obligation to be precise.

This is one of the clearest examples of responsible popularization. A writer can make the risk memorable without pretending the analogy is literal. The right posture is almost legal in its care but not legalistic in prose: comparable harm, not identical exposure; useful visualization, not medical diagnosis; PM2.5-focused method, not all pollution collapsed into tobacco.

This can become emotionally strong without becoming manipulative. It does not need to describe imaginary sick children, invented smog scenes or melodramatic dialogue. The facts are enough. The unit is enough. A city’s air translated into cigarettes already carries the shock. Adding fake drama would be both unnecessary and cheap.

The same discipline applies to carbon equivalents. A flight comparison can help readers grasp emissions, but it can also flatten complicated assumptions. Cabin class matters. Aircraft type matters. Distance matters. Load factors matter. Non-CO2 effects may matter. The more exact the claim, the more careful the method must be. If flight equivalence appears only as a supporting example of translation, it can avoid fragile current numbers. If later editions include a chart, that chart must be sourced and dated.

There is an important structural reason Chapter 5 follows freight. Shipping showed physical systems that are hidden from consumers. Air shows invisible consequences that are hidden even from direct experience. In both cases, the index

acts as an interface. Freight turns the movement of goods into rates. Cigarette equivalence turns exposure risk into a bodily metaphor.

The word interface is useful, but it should not take over. The reader does not need software vocabulary. They need the idea: a measure is something humans put between themselves and reality so they can interact with reality at all.

Sometimes the interface is a price. Sometimes it is a ranking. Sometimes it is a burger. Sometimes it is a cigarette.

Once the reader sees that, the health chapter is no longer a detour. It is the thesis under pressure.

This also provides a good place to discuss units. Units are not neutral. They decide what sort of imagination the reader will use. Micrograms per cubic meter invite technical interpretation. AQI invites color-coded public advice. Cigarettes invite bodily alarm. CO₂-equivalent invites comparison across gases and activities. Each unit carries a different emotional and analytical payload.

A unit is a small story about what matters.

That line may be worth keeping.

It explains why strange indices are not just data tricks. They are narrative decisions. The unit decides whether the reader thinks like a shopper, worker, patient, traveler, trader, engineer or citizen. A burger makes the reader a buyer. Work-time makes them a laborer. Cigarettes make them a body. Shipping makes them part of a supply chain. AI benchmarks make them a judge of machine performance.

This is why the final chapter's checklist must include the question: what unit did the index choose, and what did that choice make me feel?

Chapter 5 should leave the reader with respect for vividness and fear of it. The cigarette-equivalent measure is too useful to dismiss and too dangerous to repeat lazily. It is the perfect strange indicator: brilliant, sticky, limited and capable of becoming wrong in public at high speed.

This also helps the reader understand why public communication often chooses imperfect metaphors on purpose.

A health agency, journalist, researcher or civic tool may know perfectly well that the analogy is incomplete. They may use it anyway because the alternative is public indifference. This is a harsh bargain. If the message is too technical, it may be ignored. If it is too vivid, it may distort. The communicator stands between boredom and panic, trying to make the reader care accurately.

That problem is not unique to pollution. It appears whenever a technical risk must be made public: flood maps, heat indexes, fire danger ratings, pandemic dashboards, radiation exposure comparisons, food-safety scores. The public needs translation. Translation creates compression. Compression creates loss.

The cigarette-equivalent measure is unusually exposed because the analogy is so loaded. A cigarette is not a neutral unit like a kilometer. It is a cultural object with moral force. That force gives the index attention, and attention is precisely why the analogy must be fenced.

There is a writing lesson here. The stronger the metaphor, the stronger the caveat must be. But the caveat should not arrive as a dead bureaucratic disclaimer after the interesting part. It should be woven into the explanation. The reader can understand the analogy and its boundary in the same breath.

For example, the prose can say: the comparison does not mean the chemistry is identical; it means researchers are translating a health-risk relationship into a unit most people already understand. That sentence preserves the shock while refusing the literalism.

This is the kind of sentence a responsible popular explanation needs. It lets the book be vivid without becoming irresponsible.

Current pollution rankings are fragile unless the data is refreshed immediately before publication. Air quality is volatile. A city's reading can change with weather, season, fire, policy, traffic and monitoring coverage. The durable story is the method of translation, not a leaderboard snapshot.

The general rule is simple: use current numbers only when they are essential and source-stable; otherwise explain the measure rather than freeze a stale result.

The reader does not need to know which city was worst on a particular afternoon. They need to understand why turning air into cigarettes is both illuminating and hazardous. That idea will still matter after the afternoon changes.

Fear needs restraint here.

Fear is an efficient communicator. It is also a messy one. A cigarette-equivalent measure can make the reader care quickly because the analogy carries fear. But fear can narrow thinking. It can make people share before understanding. It can turn an analogy into accusation. In good popular nonfiction, fear is only the first spark; method and boundary must follow quickly.

The most responsible structure is: shock, explanation, limitation, transferable lesson.

Shock: the air can be translated into cigarettes.

Explanation: the comparison is based on health-risk relationships around fine particulate exposure.

Limitation: the exposure is not chemically identical to smoking, and air-quality measures are not interchangeable.

Transferable lesson: vivid equivalence can make abstraction readable, but the unit chosen shapes both understanding and misunderstanding.

This makes Chapter 5 one of the book's most important trust tests. If it works, the reader trusts the book. If it overplays the analogy, the reader has reason to distrust everything else. Health-adjacent comparisons require a higher standard because the cost of misinterpretation is higher than with burgers or lipstick.

Limitations are not a concession at the end. They are part of the beauty of the measure. A cigarette-equivalent number is interesting because it sits on a knife edge. Too technical, and nobody feels it. Too literal, and it lies. The measure's usefulness depends on balancing force and honesty.

That balance is what the book is really teaching.

The material also gives the book a chance to discuss moral weight.

Not all measurement mistakes have equal consequences. If someone overstates the Big Mac Index, they may misunderstand currency valuation. If someone overstates the Lipstick Index, they may repeat a silly recession story. If someone overstates a pollution-health analogy, they may distort risk perception. The subject has moved from amusing error to public-health communication.

This does not mean the prose must become solemn. It means the jokes need sharper boundaries. The cigarette comparison is already dramatic. It does not need a comedian standing on top of it. The humor, if any, should come from the human difficulty of understanding invisible risk, not from the risk itself.

There is also an opportunity to show respect for official units. Popular writing often treats technical language as the enemy. But micrograms per cubic meter, AQI categories and exposure windows exist for reasons. They are not bad because they are hard to feel. They are bad only when used alone with audiences who need intuitive scale.

The best communication therefore layers units rather than replacing them. It may begin with the cigarette analogy, then return to PM_{2.5} and exposure. It may use a flight comparison, then return to emissions assumptions. The strange unit opens the door; the technical unit keeps the room honest.

The pattern holds throughout: familiar unit first, precise unit nearby, limitation built in.

The problem is that equivalence can become theater. A comparison may be accurate enough for intuition and too crude for decision. It can frame guilt without explaining alternatives. It can make a number feel morally complete when the actual system involves energy mix, distance, occupancy, cabin class, non-CO₂ effects, time horizons and tradeoffs.

This is where strange measures become ethically delicate. A burger comparison can mislead your view of currencies. A cigarette comparison can mislead your sense of health risk. A

carbon comparison can mislead your sense of climate impact. The stakes are different.

That does not mean we should abandon vivid measures. It means we should be more demanding of them.

A good risk-equivalence measure needs at least four things:

First, a clear underlying pollutant or risk variable. PM_{2.5} is not “air” in general. CO₂-equivalent is not “environmental harm” in general.

Second, a clear time frame. One hour, one day, one year and one lifetime are different animals.

Third, a clear analogy boundary. The comparison must say what is being compared and what is not.

Fourth, a warning against literalism. The analogy exists to help the reader think, not to replace the underlying science.

The cigarette-equivalent index is one of the strongest examples in the book because it sits at the edge of usefulness. It is unforgettable. It is also easy to abuse. It shows exactly why strange indicators need adult supervision.

This also explains why official language often loses public attention. “PM_{2.5} concentration exceeded recommended guideline levels” may be accurate, but it does not hit the nervous system. “That air was like smoking” does. The first phrase informs. The second phrase lands.

The trick is to land without lying.

When done well, equivalence is not a gimmick. It is a bridge between measurement and consequence. It says: this invisible thing has a scale you already understand. It gives the reader enough intuitive force to care, then enough caveat to stay honest.

That is its place in the larger argument. We began with a burger because familiar objects can explain money. We moved to work-time because prices become more human when translated into life. We visited lipstick and underwear because folklore shows how seductive weak measures can be. We went to shipping because physical systems produce grounded signals. Now air shows the power and risk of trans-

lating invisible harm.

Once a measure becomes that memorable, something else happens.

People start using it.

And once people use a measure for decisions, rewards, rankings, blame or status, the measure is no longer just observing the world.

It is changing it.

Chapter 6 — When Measures Fight Back

The most dangerous moment in the life of a measure is when people begin caring about it.

That sounds unfair. A measure exists to be used. If nobody cares, it is just a lonely number in a report. But the moment a measure becomes a target, people adapt. They optimize. They polish. They route around it. They learn the rules of the scoreboard and then, with touching human consistency, begin playing the scoreboard rather than the game.

This is the broad lesson associated with Goodhart's law. In its monetary-policy origins, Charles Goodhart's point concerned the tendency of statistical regularities to break down when used for control. Later writers popularized the sharper formulation that when a measure becomes a target, it stops being a good measure. [note 20]

This is not a mystical curse. It is incentives.

A school judged mainly by test scores begins to produce test-score behavior. A company judged mainly by response time may produce fast and useless responses. A city chasing a ranking may choose visible reforms over important ones. A model lab chasing a benchmark may become excellent at the benchmark and less obviously excellent at the messy job people think the benchmark represents.

The measure has fought back.

Not always through fraud. That matters. Target-chasing exists on a spectrum. At one end is real improvement: a team sees a metric and improves the thing being measured. In the middle is narrowing: people improve the measured slice while neglecting the unmeasured rest. At the darker end is gaming, manipulation or outright fraud.

LIBOR belongs near the dark end.

For years, LIBOR was one of the most important numbers in finance. It stood for London Interbank Offered Rate, but the full name matters less than the role. LIBOR was used as a reference rate for an enormous range of financial products: loans, derivatives, mortgages and contracts whose users did not wake up each morning thinking about benchmark governance. A small movement in the benchmark could matter because so much money was priced from it.

The benchmark depended on submissions. Banks contributed estimates of borrowing rates; the submissions were used to calculate a published rate. That made LIBOR look like a technical plumbing fixture. In reality, it was also a place where human incentives could enter. Traders with positions linked to the rate had reasons to want it nudged. Submitters had discretion. The number was trusted because it looked institutional, repeated and precise. That trust was exactly what made manipulation valuable.

This is the part that makes LIBOR such a clean example. The manipulation did not require changing the weather, building a fake skyscraper or convincing millions of shoppers to buy the wrong lipstick. It required influencing a contribution to a benchmark. The human action was small compared with the economic surface area of the number. That asymmetry is the danger.

A benchmark can look like infrastructure while still depending on behavior. The public sees the published rate. Market participants see contracts. Traders see exposure. Submitters see a process. Regulators see governance. The same number sits at the intersection of all those views. If

one group can benefit from a tiny movement that another group treats as neutral plumbing, the benchmark has become a pressure point.

The phrase “critical global benchmark” sounds like legal wallpaper until you translate it. It means the number was not merely descriptive. It was wired into other things. A small movement could affect payments, valuations and trading positions. That is why benchmark manipulation is not just a technical scandal. It is a failure of trust in a measuring device people had built into the financial machine.

This is what makes LIBOR different from folklore. Nobody needed LIBOR to be charming. Nobody repeated it because it was a good dinner-party story. It mattered because contracts and institutions made it matter. The lesson is not about cultural stickiness. It is about embeddedness.

An embedded measure does not need to be memorable to be powerful.

That point belongs in a book full of vivid objects because it changes the scale of the argument. Burgers and lipstick show how measures travel through attention. LIBOR shows how measures travel through infrastructure. One wins the mind. The other moves money.

U.S. Department of Justice materials describe convictions and enforcement around schemes to manipulate LIBOR, calling it a critical global benchmark and emphasizing its impact on trillions of dollars in financial products. [note 21] The important detail for this book is not the full legal timeline. It is the mechanism: a benchmark designed to summarize market information became consequential enough that people had incentives to influence the inputs.

That is the measure-to-target problem in its most expensive clothing.

The story is not “numbers are bad.” LIBOR existed because markets needed reference points. The story is that a reference point embedded in contracts, trading positions and institutional trust needs governance equal to its power. Once money attaches to the number, the number is no longer just

a mirror. It is a lever.

Doing Business shows a different kind of pressure.

The World Bank's Doing Business report ranked economies on aspects of business regulation and enforcement. Rankings like that matter because governments care about reputation. Investors read them. Reform teams cite them. Ministers enjoy climbing them. A country moving up a table can turn a methodological score into a press release, a reform narrative and sometimes a political trophy.

The World Bank discontinued Doing Business after data irregularities in the 2018 and 2020 reports were reported internally, the next report was paused, and reviews and audits were initiated. [note 22] Again, the book does not need to become an institutional autopsy. The lesson is that rankings create status, and status creates pressure.

A ranking is not a neutral object once governments want to move inside it.

Doing Business was not a silly index. It attempted to measure real features of business regulation. That is exactly why the case is useful. The problem with consequential measures is not that they begin as jokes. Many begin as serious attempts to organize complexity. Then attention arrives. Then incentives arrive. Then the measure needs protection from the world it helped create.

Doing Business also shows a gentler-looking form of power. A ranking does not have to move money directly through contracts to change behavior. It can move attention. It can change the language of reform. It can make governments care about the parts of business regulation that appear in the score. That may produce real improvement. It may also produce reform theater: changes designed to satisfy the index more than the underlying society.

This is why rankings feel cleaner than they are. A table looks neutral. It arranges countries, firms, universities or products in order. It does not shout. It does not bribe. It does not threaten. It simply places one name above another and lets human vanity do the rest.

The power is in the ordering.

If a country moves from 76th to 52nd, the movement becomes a story even before anyone asks what changed. If another falls, the fall becomes a story too. The ranking creates a public language of improvement and embarrassment. That language is useful when the ranking is robust. It is dangerous when the ranking becomes the reform target rather than a partial view of reform.

This is not a problem only for governments. Any ranked system creates its own weather. Universities chase university rankings. Hospitals chase ratings. Apps chase store positions. Restaurants chase review averages. People chase follower counts. The mechanics vary, but the social pressure is familiar: move up, do not fall, explain the number.

Doing Business matters because the scale was institutional and global. The ranking condensed a country's regulatory environment into something comparable. That made it useful for discussion and irresistible for competition. But once countries care about the score, the score is no longer merely observation. It is part of the policy environment.

The seduction is obvious. A ranking offers a ladder. Ladders are irresistible to institutions. A government can say it moved up. A minister can point to a number. Consultants can sell advice. Journalists can simplify. Investors can compare. The ranking turns messy administrative reality into a table with winners and losers.

The problem is not comparison itself. Comparison is useful. The problem is what comparison does to the compared. Once a state knows which doors, forms, taxes, registrations and procedures are counted, the measured surface becomes special. The invisible parts remain ordinary. The ranking may encourage reform, but it also defines what reform looks like to the people chasing it.

That is why discontinuation matters to this book. The World Bank did not merely correct a typo. It stopped a flagship measurement product after irregularities and reviews damaged confidence in the report. For a measure, that is

close to institutional death. A ranking whose authority depends on trust cannot simply shrug off data irregularities and continue as if the table still speaks neutrally.

Doing Business is a story about the status economy of measurement. LIBOR is a story about the money economy of measurement. Together they show that a measure can become consequential through different channels: contracts, reputation, policy pressure, institutional competition.

Both cases turn Goodhart from slogan into narrative.

In LIBOR, the route is submission mechanics plus trading incentives.

In Doing Business, the route is ranking prestige plus institutional pressure.

In both, a measure created to summarize reality becomes part of the reality it summarizes.

The two stories also show different emotional textures of measurement failure.

LIBOR feels like a scandal because the human image is direct: traders, submissions, financial positions, a benchmark nudged where money can follow. Doing Business feels more bureaucratic, but that does not make it less revealing. It shows how even respectable measurement can become entangled with institutional incentives, prestige and pressure.

One failure is sharp. The other is atmospheric.

Both are useful because they prevent Goodhart's law from becoming a decorative quote. The law is repeated so often that it can begin to sound like a clever poster in a bad office. Cases restore its bite. They show that "measure becomes target" is not an abstract warning. It is a mechanism that can produce convictions, discontinuations, reforms, distortions and loss of trust.

The word trust is central. A measure is a trust machine. Users trust that the number was produced by a method. They trust that comparable things are being compared. They trust that contributors are not quietly bending the input. They trust that the publisher would rather correct the measure than protect its prestige. Once that trust goes, the number

may still exist, but its authority is damaged.

This is why governance cannot be an afterthought. The more power a measure has, the more visible its production needs to be. Not because transparency solves everything. It does not. But without visible method, correction and accountability, a consequential measure asks for trust while hiding the machinery that would justify it.

There is a pattern:

1. A complex reality needs a measure.
2. The measure becomes visible.
3. The visible measure becomes valuable.
4. People optimize for the valuable measure.
5. The measure begins to drift away from the reality it was meant to represent.

The same pattern appears in smaller, legal and ordinary settings. A support team optimizes response time while customers still wait for useful answers. A platform counts engagement and discovers that outrage is wonderfully engaging. A university watches rankings and adjusts what can be counted. A manager measures productivity and gets a performance of productivity.

The small examples keep the argument from feeling remote. Most readers will not manipulate a global benchmark or redesign a World Bank ranking. Many will live inside metrics. They will be measured at work, sorted by platforms, scored by institutions, nudged by dashboards and invited to confuse what is visible with what matters.

The everyday version may be harmless or useful. Measuring response time can help a support team notice neglect. Measuring uptime can make software more reliable. Measuring infection rates can save lives. The target problem does not abolish metrics. It asks for adult supervision.

The adult questions are plain:

Is the measure close to the thing we care about?

Can people improve the measured number without improving the underlying reality?

Who has reason to game it?

What happens to unmeasured parts of the system?

How will we know when the measure has gone stale?

Those questions are less memorable than “when a measure becomes a target,” but they are more useful. Slogans help people remember. Cases teach them what to do with the memory.

Nobody has to announce corruption. The metric invites behavior. Behavior adapts.

This is the mature version of the folklore problem. Lipstick and underwear survive because people love a story. Manipulated benchmarks survive, until they don’t, because people love advantage.

The more consequential the index, the more important its governance.

Governance is a boring word for a beautiful problem: who gets to decide what the number means? Who collects the data? Who checks it? Who can influence it? Who benefits if it moves? Who is harmed if it is wrong? Who notices when the measured reality has changed but the index has not?

These questions separate a toy comparison from a serious benchmark. A cigarette-equivalent pollution number may be excellent for communication but dangerous if treated as a medical calculator. A lipstick recession story may be fun in a column but silly as an investment model. A financial benchmark may move actual money. The moral and operational burden changes with consequence.

Goodhart’s law does not mean measurement is doomed. That would be lazy pessimism. It means measurement is social. A number lives among incentives, institutions and people with quarterly targets. Once a measure enters that world, it needs skepticism not because numbers are weak, but because people are inventive.

There is a practical reader payoff here. Once someone understands target effects, they become better at reading institutional life. The informed reader no longer asks only whether the number went up. They ask what behavior the

number invited.

That question travels well. It applies to benchmark rates, business rankings, school scores, product metrics and AI leaderboards. It is simple enough to remember and hard enough to answer.

There is one final distinction worth carrying forward: gaming is not always visible from the outside.

A folklore indicator can often be challenged by asking for evidence. A manipulated benchmark may look perfectly tidy in the published output. The trouble lives upstream, in submissions, incentives, methodology, sampling, governance or pressure. The public number can appear calm while the process underneath is compromised.

That is why source hierarchy is not only about authority. It is about process visibility. A first-party methodology matters because it tells the reader how the number was made. Legal or institutional records matter because they show where the process failed. A news summary may tell the story, but the underlying record gives the story weight.

The reader does not need to become a regulator. They need to notice when a number has become important enough to deserve regulation-like questions.

That awareness changes how the earlier strange measures feel. A burger index may need caveats. A financial benchmark needs controls. A lipstick story needs evidence boundaries. A government ranking needs institutional safeguards. A coding benchmark needs contamination checks. The same broad discipline applies, but the weight changes with consequence.

This is the measure reader's mature instinct: not all numbers deserve the same level of suspicion, and not all numbers can survive the same amount of pressure.

A toy indicator can be a toy. A public analogy must be careful. A decision benchmark must be governed.

The trouble begins when we treat one as another.

The burger is not just a burger; it is a measure with data choices. The work-hour comparison is not just humane; it

depends on wage definitions. The lipstick story is not just funny; it reveals our weakness for psychologically satisfying evidence. The shipping benchmark is not just grounded; it has methodology, panels, routes and market participants. The cigarette analogy is not just vivid; it risks becoming literal in public imagination.

The measure never arrives alone. It brings incentives, assumptions and failure modes.

That is why artificial intelligence belongs in this book even though it leaves conventional economics. AI appears to be a new domain, full of new vocabulary and new machines. But its benchmark culture is an old problem in a new suit.

The question is familiar: how do we measure something too complicated to see directly?

The answer, as usual, is: badly, usefully and with consequences.

Chapter 7 — Benchmarks for Machines

Artificial intelligence did not invent the benchmark problem.

It merely made it faster, louder and more expensive.

The AI world is full of measures: model scores, benchmark suites, leaderboards, capability reports, readiness rankings, safety evaluations, coding tests, math tests, reasoning tests, multimodal tests, agent tests and charts showing lines climbing upward with the calm menace of a species that has discovered PowerPoint. The names are new. The problem is old.

We are trying to measure something too large and slippery to observe directly.

The Stanford AI Index is one attempt to organize the field. Stanford HAI describes the report as tracking, collating, distilling and visualizing AI-related data for policymakers, researchers, executives, journalists and the public. [note 23] That is index work in the broadest sense: take a messy do-

main and give people a structured way to read it.

SWE-bench is narrower and more concrete. It evaluates language models on real-world GitHub issues by giving a model a codebase and an issue, then asking it to generate a patch. The official project describes it as a benchmark for real-world software issues collected from GitHub, where a model must generate a patch that resolves the described problem. [note 24]

This is a better test than asking a model to produce a tidy programming answer in isolation. Real software work involves existing code, tests, constraints and the unpleasant fact that changing one line can annoy six others.

It is also a perfect example of benchmark lifecycle.

SWE-bench began by making evaluation more realistic. The original paper introduced a set of software-engineering problems drawn from GitHub issues and pull requests across popular Python repositories. The benchmark was hard enough that early model performance was low; that difficulty made it respected. It seemed to move AI evaluation closer to real work: not “write a function from scratch,” but “enter this existing mess and fix the issue without breaking the house.” [note 25]

That move mattered because many earlier coding evaluations were too clean. They asked for answers in the abstract: produce a function, solve a programming puzzle, complete a snippet. Useful, but not the same as software maintenance. Real repositories are rude. They contain conventions nobody explains twice, tests that encode old decisions, dependencies, naming habits, hidden coupling and comments written by someone who either knew too much or not enough. The benchmark’s appeal came from dragging models into that world.

The unit was no longer “can the model write code?” It became “can the model resolve this issue in this existing repository?”

That is a much better unit.

It is also a more expensive one. To solve such a task, a

system may need retrieval, long context, planning, editing, testing and iteration. Once those surrounding tools matter, the benchmark no longer measures a bare language model in isolation. It measures a model inside an evaluation setup. That may be the right thing to measure if the product is also a system. But the public headline often forgets the distinction.

A score can belong to a model, a scaffold, a toolchain, a prompt strategy, a test harness or all of them at once. If the headline says only “Model X scored Y,” some of the machinery has disappeared.

Then the benchmark became important.

Importance changes a benchmark. Once researchers, labs, users and journalists pay attention, performance becomes reputation. Systems are tuned, harnesses improve, scaffolds appear, test-time strategies spread, and the benchmark starts measuring not only raw model ability but the entire evaluation ecosystem around the model. This may be legitimate progress. It may also mean the score is no longer as simple as it looks.

The next step was refinement. OpenAI released SWE-bench Verified as a human-validated subset intended to evaluate the ability to solve real-world software issues more reliably. [note 26] That is a good benchmark story: the first measure exposes something useful; then people inspect the measure itself and create a cleaner subset.

The need for a verified subset is itself revealing. It says the original benchmark was useful enough to matter and imperfect enough to improve. Some tasks in a broad benchmark may be ambiguous, broken, underspecified, too easy, too hard, or poorly aligned with what the benchmark claims to test. Human validation tries to clean the signal.

This is not a scandal. It is measurement maturing.

Every serious measure goes through this if it survives long enough. The first version proves the category. The second version removes noise. Later versions fight saturation, contamination and strategic adaptation. A benchmark that never changes may be stable because it is excellent. It may

also be stable because nobody important is testing it.

SWE-bench Verified therefore gives the book a concrete lifecycle rather than a generic warning. Benchmark created. Benchmark respected. Benchmark inspected. Benchmark refined. Benchmark optimized against. Benchmark later questioned for frontier measurement.

But the lifecycle did not end there. Later, OpenAI argued that SWE-bench Verified was increasingly contaminated and mismeasuring frontier coding progress, recommending SWE-bench Pro for that frontier use case. [note 27] Whether one accepts every conclusion or not, the measurement lesson is excellent. A benchmark can be valuable, then become saturated or contaminated for the very systems it helped improve.

Contamination is a particularly modern version of an old problem. In a normal school test, students may practice past papers. In AI, training data, public solutions, benchmark discussion and scaffold design can all blur the line between solving the task and having absorbed its shape. The benchmark becomes part of the environment. The more famous it becomes, the harder it is to keep it pristine.

That is one reason AI evaluation feels so unstable from the outside. The test is not simply administered to passive systems. The test becomes part of the culture building the systems. Researchers read it. Engineers optimize around it. Products advertise against it. Users quote it. Datasets circulate. Harnesses are replicated. The benchmark becomes famous, then fame becomes a variable inside the benchmark's own usefulness.

This is measurement folding back on itself.

The same thing happens in slower fields, but AI makes it visible because the loop is compressed. A school exam may take years to become overcoached. A university ranking may take years to reshape institutional behavior. An AI benchmark can become a target almost immediately because the field is built around public comparison and rapid iteration.

This does not make the benchmark dishonest. It makes it

perishable.

Saturation is the happier-looking problem. A benchmark can become too easy for frontier systems. That sounds like success, and in one sense it is. A hard test forced progress. Then progress made the test less discriminating. The measure did its job and then lost sharpness at the top end.

Both problems show why a benchmark is not a monument. It is a tool with a service life.

The phrase “service life” is important because it avoids moral panic. A benchmark that becomes saturated has not necessarily failed. It may have succeeded so well that it needs replacement. A benchmark that becomes contaminated has not necessarily been useless. It may have become too visible to serve its original purpose cleanly.

That is a calmer way to read benchmark turnover. New tests do not always prove old tests were bad. Sometimes they prove the field moved. Sometimes they prove the old test was too narrow. Sometimes they prove the incentives became too strong. Often they prove several things at once.

It is worth resisting the easy reaction: “If benchmarks keep changing, none of them mean anything.” That is too simple. Road maps change and still help travelers. Medical tests improve and still diagnose. Economic indices revise methods and still inform. The question is not whether a benchmark changes. The question is whether the change is explained, justified and matched to the claim being made.

This is where source notes, methodology pages and benchmark papers matter. They are not decorative paperwork. They are the benchmark telling the reader what sort of trust it wants.

That phrase belongs in AI evaluation, but it also belongs in the rest of the book. A consumer index can become stale. A shipping benchmark can need methodological governance. A readiness ranking can become politicized. A health analogy can travel beyond its assumptions. A measure has a life cycle: invention, adoption, misuse, refinement, replacement or decline.

SWE-bench simply makes that cycle easier to see because it is compressed into a fast-moving technical field.

This is Goodhart with a GPU budget.

A benchmark is created. It is difficult. Everyone respects it. Then models improve, tools adapt, prompt strategies spread, evaluation harnesses become familiar, and the benchmark ages. A new benchmark appears to fix the old benchmark's weaknesses. It too becomes famous. Then it becomes a target. The frontier moves. The scoreboard grows teeth.

This is not unique to AI. Education tests, financial benchmarks, university rankings and business-climate indices have all been through versions of the same cycle. AI compresses the cycle because the field moves quickly and benchmark performance is commercially potent. A few percentage points on a public leaderboard can influence reputation, investment, procurement and user belief. That is a large burden for a test.

SWE-bench is useful precisely because it tries to move evaluation closer to real work. But "closer" is not "identical." Professional software engineering includes ambiguous requirements, team coordination, maintenance risk, security judgment, architecture taste, product context, deployment discipline and the social art of not turning every bug fix into a new platform. A benchmark can capture some of that. It cannot capture all of it.

This distinction is not anti-benchmark. It is pro-benchmark sanity.

Without coding benchmarks, the field would drown in anecdotes. One person would say a model saved their week. Another would say it destroyed their afternoon. Both might be true. Benchmarks create a shared object of argument. They let researchers and builders compare systems under known conditions. They make progress legible.

But legibility is not completeness.

The benchmark tells us what happened inside the harness. It does not tell us whether the system understands a company's architecture, respects security constraints, nego-

tiates product tradeoffs, or knows when not to touch a file. Those capabilities may be tested elsewhere. They may remain untested. Either way, one score should not be allowed to smuggle them in.

This is the same mistake people make with all strong proxies. The better the proxy, the more tempting it becomes to generalize. A weak proxy invites skepticism. A strong proxy invites overreach.

This is not an insult to SWE-bench. It is what makes a benchmark readable. You respect it by keeping its scope intact.

The AI Index has a different failure mode. It is not one task benchmark but a broad report trying to organize a sprawling field: research, technical performance, responsible AI, economy, science, medicine, education, policy and public opinion. That kind of index can be comprehensive and still incomplete. It can help readers orient themselves without digesting the field on their behalf.

Government AI readiness rankings have yet another failure mode. They take institutions, infrastructure, data, skills and policy conditions and compress them into rankings. Policymakers need that compression. Journalists need it. Investors and consultants certainly enjoy it. But a country is not “ready” in one dimension. A country may have strong technical talent and weak governance. It may have policy ambition and poor data infrastructure. It may have money and low public trust. A ranking can begin the conversation. It should not end it. [note 28]

AI also introduces the problem of demo-as-index. Demonstrations are often more persuasive than evaluations. A model solving a flashy task on video can move public belief more than a careful report. The demo becomes a strange indicator of its own: if it can do this one vivid thing, perhaps it can do the whole category. Sometimes that inference is fair. Often it is a magic trick performed by editing, selection and human patience.

A useful distinction is simple:

Capability: what the system can do under specified conditions.

Reliability: how often it does it correctly across messy cases.

Usefulness: whether the result matters in the real workflow.

Benchmarks often measure capability better than reliability, and reliability better than usefulness. The gaps are where disappointment lives.

This three-part distinction gives the argument its practical spine.

A model can demonstrate capability by solving a benchmark task. That proves something. It may be an impressive something.

Reliability asks whether the performance holds across many messy cases, with varied repositories, changing dependencies, incomplete issues and no hidden hand-holding.

Usefulness asks whether the result fits into actual work. Did it save time? Did it create review burden? Did it introduce subtle risk? Did it preserve project intent? Did the human still need to understand everything before trusting it?

Most public benchmark talk compresses all three into one score. That compression is useful for headlines and dangerous for decisions.

The reader does not need to know the technical details of every benchmark to benefit from this distinction. It applies whenever a system is evaluated. A school exam measures capability under exam conditions. A driving test measures enough performance to issue a license, not a lifetime of safe driving. A product demo measures a selected flow. A business ranking measures selected regulatory conditions. A software benchmark measures selected tasks under selected assumptions.

The phrase “under selected assumptions” is not a footnote. It is the whole game.

Consider what happens when a model benchmark score enters a business decision. A buyer may not know the har-

ness details. They see a chart. Procurement sees a leader. A founder sees a claim. A journalist sees a headline. The score has moved from evaluation into persuasion. If the benchmark task matches the buyer's real workflow, the score may help. If it does not, the score becomes expensive theater.

This is the same movement from Chapter 6 in technical clothing. A measure becomes important, then begins shaping behavior around it. Labs optimize. Buyers compare. Investors notice. The benchmark participates in the market it claims to describe.

This is why AI belongs after Goodhart rather than before it. The reader has already learned to ask what happens when people care about a number. AI benchmarks answer: they build around it at speed.

The tone can be skeptical without being anti-progress. Some model improvements are real. Benchmarks help reveal them. A coding model that solves more real repository issues is not doing nothing. But a benchmark score is the beginning of evaluation, not the end of adoption. For a serious user, the next questions are workflow fit, reliability, review burden, security, cost and failure mode.

Those questions are less glamorous than a leaderboard. They are also where usefulness lives.

Current-score obsession is deliberately avoided for a reason. Exact model rankings may be stale by publication. The durable story is the mechanics: a benchmark frames a task, public scores travel, optimization pressure appears, and scope matters. The reader does not need the current number-one model to understand why "number one" is never a complete sentence.

AI belongs near the end rather than the beginning because it proves the book's pattern in a modern domain. If the book opened with AI, it would risk becoming a technology book. Here, AI becomes the latest example of an old measurement problem, not the main event.

The machines may be new. Our weakness for leaderboards is ancient.

When an AI benchmark score appears, ask what task it measures. If it is a coding benchmark, ask whether the task resembles real coding or a contest. Ask what tools were allowed, how failures were counted, whether tests were reliable, whether data leakage is plausible and whether the score reflects a model, a scaffold or a whole evaluation stack.

The same logic applies to all strange indicators. A burger measures one kind of purchasing-power story. Work-time measures affordability through wages. Lipstick measures, at best, a consumer-behavior hypothesis. Shipping measures physical transport markets. Cigarette equivalents measure risk translation. AI benchmarks measure task performance under a defined harness.

Each is useful when kept in its cage.

Each becomes dangerous when released into general conversation without a tag.

As systems become more complex, the demand for simple indicators rises. Nobody wants to read every model card, audit report, evaluation detail and methodology note before forming a view. People want a number. They want a rank. They want a color, badge, score or object. This is not a failure of intelligence. It is a survival mechanism.

But the more we need compression, the more we need to inspect the compressor.

That brings the book to its final task. We have walked through burgers, phones, lipstick, underwear, ships, air, benchmarks and manipulated targets. The question is no longer whether strange indicators are serious. They clearly can be. Nor is it whether they are dangerous. They clearly can be that too.

The question is how to read them without becoming either a bore or a sucker.

Chapter 8 — How to Read a Weird Index

A weird index should be enjoyed before it is trusted.

That order matters. Suspicion without enjoyment becomes sterile; it misses why the measure worked in the first place. Enjoyment without suspicion is worse; it lets myths enter the mind, sit down and become furniture. The useful reader does both. Let the indicator land. Then inspect the landing gear.

By now the pattern should feel familiar. A burger can make purchasing power visible. Work-time can make price personal. Lipstick can turn recession anxiety into a story people want to repeat. Freight rates can show the physical economy moving underneath consumer life. Cigarette equivalents can turn invisible risk into a bodily metaphor. LIBOR, Doing Business and AI benchmarks can show what happens when measures become consequential.

The final skill is not cynicism. Cynicism would say all indices are lies. That is false and lazy. Many indices are useful. Some are essential. The modern world cannot function without prices, benchmarks, rankings, ratings, standards and indicators. The better skill is conditional respect.

Respect the attempt to measure.

Interrogate the translation.

Watch the incentives.

Keep the object in proportion.

Here is a practical way to do that.

First: what is the object?

Every good strange indicator has an object or unit that carries the idea. A Big Mac. An hour of work. A lipstick. A container. A cigarette-equivalent. A benchmark score. If the object is vague, the measure is probably vague too. “Confidence,” “innovation,” “quality” and “readiness” can be measured, but only after someone makes a series of choices. Those choices are where the story lives.

Second: what is the source?

A first-party methodology is better than a viral chart. An official guideline is better than a recycled blog post. A serious paper is better than a phrase repeated by people who enjoy

sounding interesting. This does not mean only institutions can be right. It means source quality matters most when the claim is most memorable.

Third: what is the method?

How is the number produced? Is it a price observation, a survey, a model, an expert assessment, a scraped dataset, a benchmark harness, a ranking formula or a journalist's clever comparison? The method does not need to be perfect. It needs to be visible enough to criticize.

Fourth: what is the scope?

The Big Mac Index works where Big Macs and data exist. A work-time index depends on selected wages and products. The Baltic Dry Index concerns dry bulk shipping, not all trade. SWE-bench concerns specific software tasks, not all intelligence. A measure that clearly states its scope deserves more trust than one that quietly expands into everything.

Fifth: what is the strength of the claim?

This is the lesson from lipstick, underwear and skyscrapers. "This behavior sometimes happens" is not the same as "this diagnoses recession." "This diagnoses mood" is not the same as "this predicts the future." A weak claim can become false by promotion. Before believing the sentence, identify how much weight it is trying to carry.

Sixth: what changes when people optimize for it?

This is the Goodhart question. If the measure becomes important, who can game it? Who can polish it? Who can redefine it? Who benefits? A benchmark attached to money, status or policy is never just a passive description. It becomes part of the system it measures.

Seventh: what would prove it wrong?

A measure that cannot be embarrassed is not a measure. It is a belief with formatting. If lipstick sales fall during a downturn, what happens to the Lipstick Index? If a benchmark score rises while real-world performance disappoints, what happens to the benchmark story? If a readiness ranking fails to predict readiness, does anyone notice?

This checklist is deliberately plain because the skill of read-

ing weird indices is less about advanced mathematics than intellectual hygiene.

The checklist also works because it slows down the moment of belief.

Most bad index reading happens too quickly. The reader sees a clever unit, a sharp ranking or a surprising comparison and immediately turns it into a conclusion. The pause is where judgment lives. Not a long pause. Not a dissertation. Just enough time to ask whether the measure has earned the sentence being built on top of it.

Take the burger. The object is clear. The source is identifiable. The method is explainable. The scope is limited. The claim is modest when handled properly: useful purchasing-power intuition, not currency destiny. That is a healthy strange index.

Take lipstick. The object is clear and the story is memorable. The source trail is weaker, the evidence is contested and the claim often expands beyond what it can support. That does not make the story worthless. It means the right use is cultural and diagnostic, not predictive.

Take LIBOR. The object is not charming at all, but the source and consequence are clear. It shows that a benchmark can become so embedded in contracts and markets that manipulation of the measure becomes manipulation of the world built around it.

Take SWE-bench. The object is a benchmark task: real repository, real issue, generated patch. The source and method are visible enough to inspect. The scope is narrower than public excitement often suggests. The lifecycle — benchmark, verified subset, saturation and contamination concerns — shows that even good measures age.

The same questions produce different verdicts. That is the point. A good reader does not apply one mood to every index. They do not either admire everything or debunk everything. They classify.

Some measures deserve trust.

Some deserve use with caveats.

Some deserve laughter and quarantine.

Some deserve governance.

Some deserve retirement.

This is a better outcome than a single grand theory. The world does not need another person declaring that all metrics are fake or all data is truth. It needs readers who can tell a useful proxy from a shiny trap.

Do not confuse vivid with true.

Do not confuse official with complete.

Do not confuse measured with important.

Do not confuse ranked with understood.

Do not confuse a proxy with the thing.

These rules are simple enough to sound obvious. They are not obvious in use.

Markets, media and institutions constantly reward over-reading. A ranking becomes a verdict because verdicts travel. A benchmark becomes a claim about intelligence because intelligence sells. A consumer anecdote becomes a recession signal because recession anxiety needs objects. A health analogy becomes literal because literal fear spreads faster than careful explanation.

The forces pushing toward exaggeration are strong. The discipline pushing back must be simple enough to survive contact with headlines.

That is why the book ends with proportion rather than certainty. Proportion is the underrated virtue in measurement. It asks the reader to keep the measure at the right size. The Big Mac Index is big enough to teach PPP and too small to run foreign-exchange strategy. Work-time is big enough to reveal affordability and too small to explain all inequality. Lipstick is big enough to reveal the appeal of recession folklore and too small to forecast the economy. SWE-bench is big enough to test a class of coding performance and too small to settle the future of software work.

Right-sized measures are useful. Inflated measures become myths.

The title objects now have their jobs. The burger is the

friendly proxy: imperfect, methodical enough and excellent at opening a serious idea. The lipstick is the seductive myth: psychologically satisfying, culturally useful and empirically unsafe when overpromoted. The underwear is the comic warning: memorable enough to travel, weak enough to handle carefully.

Together they make the promise honest. Strange indicators are not brilliant or idiotic as a class. They occupy a spectrum. The reader's job is to place each one on that spectrum before using it.

That spectrum extends beyond consumer objects. A port, a pollutant and a software benchmark can all become strange measures once they are translated into a form people can use. The names change. The reading discipline remains.

A good strange indicator says: here is one useful way to look.

A bad one says: here is reality.

The world will remain too large to see directly. We will keep needing proxies. We will keep turning abstraction into objects, hours, ships, cigarettes and scores. That is not a weakness. It is one of the ways humans think.

The discipline is to carry two thoughts at once.

The first is delight: what a strange, clever way to see the world.

The second is suspicion: what exactly disappeared in the translation?

Hold both, and the burger becomes useful. The lipstick becomes interesting. The underwear becomes funny but contained. The ship becomes a signal. The cigarette becomes a warning with boundaries. The benchmark becomes a tool instead of a shrine.

Curiosity first.

Then method.

Then limits.

Then incentives.

Then proportion.

Enjoy the measure.

Do not worship it.
That is enough.

Notes

1. The Economist. *Data and methodology for the Big Mac index*. GitHub / The Economist data repository, n.d..
github.com/TheEconomist/big-mac-data
2. Investopedia. *Understanding the Starbucks Index*. Investopedia, n.d..
www.investopedia.com/terms/s/starbucks-index.asp
3. Owner / ABVX. *Owner-supplied unusual indices research source pack*. Local source pack, n.d..
[books/source-packs/unusual-indices-book-source-pack.json](#)
4. UBS. *Prices and Earnings 2015*. UBS Prices and Earnings, 2015. files.illinoispolicy.org/wp-content/uploads/2015/09/ubs-price...
5. World Economic Forum. *Who has to work the longest to buy an iPhone?*. World Economic Forum, 2015-12-07.
www.weforum.org/stories/economic-growth/who-has-to-work-the-l...
6. Numbeo. *Cost of Living database*. Numbeo, n.d..
www.numbeo.com/cost-of-living
7. RBC / secondary market commentary. *Lipstick Index secondary market commentary*. Research lead, n.d..
8. J.P. Morgan Wealth Management / Chase. *What is the Lipstick Index?*. Chase, 2026-06-05.
www.chase.com/personal/investments/learning-and-insights/arti...
9. Secondary business / medical-economics commentary. *Unusual recession-indicator overview*. Research lead, n.d..
10. Jason Barr, Bruce Mizrach, Kusum Mundra. *Skyscraper Height and the Business Cycle: Separating Myth from Reality*. Applied Economics, 2015.
ideas.repec.org/a/taf/applec/v47y2015i2p148-160.html
11. Market Urbanism. *Skyscrapers as Economic Indicators*. Market Urbanism, 2008-08-26.

- marketurbanism.com/2008/08/26/skyscrapers-as-economic-indicators
12. Baltic Exchange. *Baltic Exchange methodology*. Baltic Exchange, n.d.. www.balticexchange.com/en/data-services/Methodology.html
 13. Freightos. *Freightos Baltic Index*. Freightos, n.d.. www.freightos.com/enterprise/terminal/freightos-baltic-index-...
 14. Ember. *Energy data*. Ember, n.d.. ember-energy.org/data
 15. Electricity Maps. *Electricity Maps*. Electricity Maps, n.d.. app.electricitymaps.com/map
 16. Berkeley Earth. *Air Pollution and Cigarette Equivalence*. Berkeley Earth, n.d.. berkeleyearth.org/air-pollution-and-cigarette-equivalence
 17. OpenAQ. *AQI Hub*. OpenAQ, n.d.. openaq.org/about/initiatives/aqi-hub
 18. World Health Organization. *WHO global air quality guidelines*. WHO, 2021. www.who.int/publications/i/item/9789240034228
 19. Our World in Data. *CO2 and greenhouse gas emissions*. Our World in Data, n.d.. ourworldindata.org/co2-and-greenhouse-gas-emissions
 20. Charles Goodhart / later economics and policy usage. *Goodhart's law / measure-target risk*. Conceptual source family, n.d.. link.springer.com/chapter/10.1007/978-1-349-17295-5_4
 21. U.S. Department of Justice. *Two Former Deutsche Bank Traders Convicted for Role in Scheme to Manipulate Critical Global Benchmark Interest Rate*. Justice Department Archives, 2018-10-17. www.justice.gov/archives/opa/pr/two-former-deutsche-bank-trad...
 22. World Bank Group. *World Bank Group to Discontinue Doing Business Report*. World Bank statement, 2021-09-16. www.worldbank.org/en/news/statement/2021/09/16/world-bank-gro...
 23. Stanford HAI. *AI Index Report*. Stanford AI Index, 2026. hai.stanford.edu/ai-index/2026-ai-index-report
 24. SWE-bench project. *SWE-bench*. GitHub / project site, n.d.. github.com/swe-bench/SWE-bench
 25. Carlos E. Jimenez et al.. *SWE-bench: Can Language*

- Models Resolve Real-world GitHub Issues?*. arXiv, 2023.
arxiv.org/abs/2310.06770
26. SWE-bench / OpenAI collaboration. *SWE-bench Verified*.
SWE-bench project site, 2024.
www.swebench.com/verified.html
27. OpenAI. *Why SWE-bench Verified no longer measures
frontier coding capabilities*. OpenAI, 2026.
[openai.com/index/why-we-no-longer-evaluate-swe-bench-
verified](https://openai.com/index/why-we-no-longer-evaluate-swe-bench-verified)
28. Oxford Insights. *Government AI Readiness Index*. Oxford
Insights, 2025. [oxfordinsights.com/ai-
readiness/government-ai-readiness-index...](https://oxfordinsights.com/ai-readiness/government-ai-readiness-index...)